

Judicial Impact Assessment (JIA): A Proactive Framework for Municipal Regulatory Resilience in the Post-*Loper Bright* Era

Oksana Manko¹

ABSTRACT

Loper Bright Enterprises v. Raimondo eliminated *Chevron* deference, requiring federal courts to exercise independent judgment when interpreting ambiguous statutes. Existing scholarship has analyzed this shift almost entirely at the federal level, treating it as a recalibration between courts and agencies. This Article argues that the decision's most consequential effects fall on a layer of government largely absent from that debate: the roughly 90,000 municipalities that draft ordinances, permits, and licensing rules under statutes as ambiguous as anything an agency interprets, with far fewer resources to defend a contested reading. Compounding circuit-to-circuit divergence, an unsettled revival of *Skidmore* deference, and the elimination of any practical statute of limitations for facial challenges under *Corner Post*, municipalities now face what this Article calls a judicial lottery: the same interpretive method, applied by different panels in good faith, regularly produces different answers to comparable questions. This Article proposes the Judicial Impact Assessment (JIA), a four-block methodology that scans regulatory text for interpretive vulnerability before adoption, scores litigation and preemption risk through the Municipal Litigation Vulnerability Index, forecasts the fiscal consequences of an adverse ruling, and translates those findings into concrete drafting responses. Comparative analysis of five foreign

¹ Candidate of Legal Sciences (Ph.D. equivalent); Researcher in Administrative and Local Governance Law. Formerly Chief Consultant in the Administration of the President of Ukraine (2011–2024), where she worked on constitutional reform, decentralization, and local self-government. This Article reflects her current work developing the Judicial Impact Assessment (JIA) framework for U.S. municipalities in the post-*Loper Bright* era. All errors are her own. oksana.manko.academic@gmail.com

systems shows that no existing framework, domestic or foreign, is built to answer this question. Retrospective application to *Sackett v. EPA* and *City and County of San Francisco v. EPA*, together with a bounded comparative lesson from Ukraine’s decentralization reform, isolates a common condition this Article calls interpretive dependence, tests the framework against real outcomes, and proposes an empirical agenda for validating it further.

Keywords: municipal law, administrative law, *Loper Bright*, *Chevron* deference, statutory interpretation, regulatory risk assessment, local government law, judicial review

I. INTRODUCTION

1.1 The End of Deference

*Chevron U.S.A. Inc. v. Natural Resources Defense Council, Inc.*² governed federal regulatory practice for forty years. Its rule was simple. If a statute was ambiguous, courts deferred to the agency’s reasonable reading of it. That rule is gone.

On June 28, 2024, the Supreme Court decided *Loper Bright Enterprises v. Raimondo*.³ The Court held that the Administrative Procedure Act requires judges to exercise independent judgment when a statute is unclear. It did not trim *Chevron*. It overruled it. Courts must now determine the best reading of a statute themselves, using traditional tools of construction — though they may still draw on an agency’s reasoning as one interpretive aid, under the older, weaker standard from *Skidmore v. Swift & Co.*⁴ An agency’s interpretation can still carry weight. It simply cannot control the outcome.

² 467 U.S. 837 (1984).

³ 603 U.S. 369 (2024).

⁴ 323 U.S. 134 (1944).

Most commentary on *Loper Bright* has stayed at the federal level. That is fair enough as far as it goes. But it misses something. *Loper Bright* is usually described as a shift between two federal branches, courts and agencies. In practice, its effects run further down, into a layer of government that gets almost no attention in this debate: the roughly 90,000 municipalities that write zoning codes, environmental permits, and licensing rules under statutes that are often just as ambiguous as anything an agency interprets — and with far fewer resources to defend their reading of them. Municipalities differ from federal agencies in a way this literature has not yet reckoned with. Federal agencies generally operate within a more coherent statutory and institutional framework than municipalities do. A city operates under an accumulation of state enabling statutes, home-rule provisions, and federal grant conditions, each drafted at a different time, by a different legislature, often without reference to the others. That layered structure is not a smaller version of the agency problem. It is a different problem, and it is why municipalities merit treatment as their own category of post-*Loper Bright* risk, not as a footnote to administrative law.

1.2 Municipal Vulnerability and the “Judicial Lottery”

Most of a city attorney’s work is unglamorous: a noise ordinance, a wastewater permit, a zoning variance. *Loper Bright* has quietly transformed that unglamorous work.

Federal agencies can build a deep administrative record to support their interpretation. Most cities cannot. A well-developed record now matters far more because it supports the kind of reasoning *Skidmore* treats as persuasive.⁵ Private parties can often simply avoid a bad legal position; cities cannot. State law and federal grant conditions force them to regulate, often using the same vague statutory language courts are now free to read however they see fit — and across a genuinely fractured judicial map. A zoning ordinance drafted around one circuit’s reading of a preemption

⁵ See *infra* Part II.2

statute can be struck down the moment a different panel, or a neighboring circuit, reads the same words differently. This Article refers to this phenomenon as the judicial lottery. It is a concrete problem: cities can no longer easily forecast whether a routine regulation will hold up, and the cost of guessing wrong keeps growing.

The uncertainty this creates is not confined to federal courtrooms. Florida offers a compact illustration: within the same permitting cycle, a federal magistrate applied *Sackett* to release a landowner from Clean Water Act jurisdiction,⁶ while months later a federal appellate ruling stripped the state’s own delegated wetlands-permitting authority back to federal agencies.⁷ Trade and practitioner commentary confirms that which regulatory regime actually governs a given wetland in a given state remains, in the industry’s own words, “somewhat murky.”⁸ A municipality sitting between these two authorities has no more clarity than the courts and agencies do.

Two cases at the Supreme Court level show what is at stake more starkly. *Sackett v. EPA*,⁹ decided the year before *Loper Bright*, had already narrowed “waters of the United States” through rigorous textual construction — foreshadowing the interpretive posture *Loper Bright* would soon make general, and unsettling years of permitting practice built on the older reading. *City and County of San Francisco v. EPA*¹⁰ went the other way: a city used strict textual argument to defeat the EPA’s own expansive reading of its permitting authority, prevailing 5–4 on the ground that Section 1311(b)(1)(C) of the Clean Water Act does not authorize “end-result” conditions untethered to

⁶ *United States v. Sharfi*, No. 2:21-cv-14205-KAM, 2024 WL 4483354 (S.D. Fla. Sept. 21, 2024) (R&R), *adopted* 2024 WL 5244351 (Dec. 30, 2024).

⁷ See Earthjustice, *Court Victory for Florida Wetlands Protection* (2026) (Eleventh Circuit affirming that EPA and FWS violated the Endangered Species Act in approving Florida’s delegated 404 wetlands-permitting program).

⁸ E.A. Crunden, *Post-Sackett, Chaos Erupts for Wetlands Oversight*, E&E News (June 2, 2023) (quoting a permitting expert describing post-*Sackett* uncertainty as “far from clear”); Frederick Aschauer & Katherine Hupp, *From Maui to San Fran: A Look at the Changing Clean Water Act Through the Eyes of SCOTUS*, Florida Specifier (July/Aug. 2025).

⁹ 598 U.S. 651 (2023).

¹⁰ 604 U.S. 334 (2025).

specific limits.¹¹ One case narrowed federal reach against a landowner; the other checked it in favor of a city. But both teach the same lesson: the party best able to anticipate how a court will read ambiguous text enters litigation with a substantial advantage. Cities are, as a class, the least prepared and most exposed litigants in that contest.

1.3 The Gap in the Literature

Two bodies of scholarship address adjacent questions, yet they rarely intersect.

Administrative law scholars have debated what *Loper Bright* will do to interpretive stability. Thomas Merrill expects it to “cause legal disagreement to proliferate” across circuits, reproducing at the judicial level the instability once attributed to agencies under *Chevron*.¹² Daniel Hemel takes the opposing view: replacing agency “flips” with judicial “splits” is, on balance, a desirable trade, since divergence across circuits can eventually be resolved through Supreme Court review in a way that instability across time cannot. Cass Sunstein sits between them, treating the decision as a genuine constitutional correction while conceding that its practical effects are still unclear; Jonathan Adler and Ilya Somin add a further caution against overstating its significance either way. None of this scholarship was written with a city hall in mind.

Municipal law scholars, for their part, have built a strong account of a different kind of vulnerability — one running horizontally, between a city and its own state, rather than vertically, between a city and a reviewing court. Richard Briffault’s account of what he calls “the new

¹¹ 604 U.S. 334 (2025) (holding that “§1311(b)(1)(C) does not authorize the EPA to include ‘end-result’ provisions in NPDES permits,” and that “[d]etermining what steps a permittee must take to ensure that water quality standards are met is the EPA’s responsibility”).

¹² Thomas W. Merrill, *The Demise of Deference — And the Rise of Delegation to Interpret?*, 138 Harv. L. Rev. 227 (2024); Daniel J. Hemel, *Flips and Splits in Administrative Law*, 74 Duke L.J. 1703 (2025); Cass R. Sunstein, *Our Marbury: Loper Bright and the Administrative State*, 74 Duke L.J. (forthcoming 2025); Jonathan H. Adler, *From “Deference” to “Respect” — The Real Import of Loper Bright*, Volokh Conspiracy (July 3, 2024); Ilya Somin, *The Supreme Court’s Decision Overruling Chevron Is Important — But Less So Than You Might Think*, Volokh Conspiracy (June 28, 2024).

preemption” documents a shift from classic, subject-matter preemption toward something more aggressive: state laws that impose personal liability or fiscal sanctions on local officials, or that strip local regulatory authority outright, closely tied to the partisan and ideological polarization between state legislatures and the cities within them.¹³ Nestor Davidson traces a related dynamic, describing how sharply rising political and cultural conflict between states and cities has turned questions of local authority into a live battleground, and arguing for a more explicit reckoning with the normative stakes on both sides.¹⁴ Richard Schragger and Erin Scharff have each examined, from different angles, how much practical and fiscal power cities retain despite their formal legal subordination to the states that charter them.¹⁵

This literature and the one just discussed are not in tension. They are answering different questions. A city with unquestioned home-rule authority to regulate a given subject — no preemption problem in Briffault’s or Davidson’s sense at all — can still face the judicial lottery this Article describes, if the enabling statute it relies on, or a federal grant condition attached to its funding, turns on language a court now reads independently rather than deferentially. The preemption literature asks who gets to decide, as between a state and its cities. This Article asks a different question: once a city has the authority to decide, how exposed is its decision to a court’s own, unpredictable reading of the text that authority rests on.

This emerging concern has begun to surface in practitioner commentary as well: one recent client alert observes that “[m]uch has been written about the seismic effect *Loper Bright* has had on

¹³ Richard Briffault, *The Challenge of the New Preemption*, 70 Stan. L. Rev. 1995, 1999–2008 (2018) (documenting “punitive” and “nuclear” preemption and their connection to state-city political polarization); Briffault et al., *Principles of Home Rule for the Twenty-First Century* (Nat’l League of Cities 2020).

¹⁴ Nestor M. Davidson, *The Dilemma of Localism in an Era of Polarization*, 128 Yale L.J. 954 (2019).

¹⁵ Richard C. Schragger, *City Power: Urban Governance in a Global Age* (2016); Erin Adele Scharff, *Powerful Cities?: Limits on Municipal Taxing Authority*, 91 Wash. U. L. Rev. 1225 (2014); cf. Robin K. Craig, *The Impact of Loper Bright v. Raimondo: An Empirical Review of the First Six Months*, 109 Minn. L. Rev. 2671 (2025) (state-court dimension, not municipal-specific).

federal administrative law,” while “[l]ess discussed is the aftershock effect the decision may yet have on state administrative law.”¹⁶ But neither the academic municipal law literature nor this early practitioner notice has yet produced a working tool that turns *Loper Bright*’s doctrine into something a city attorney can use while drafting an ordinance.

Existing practitioner tools do not close this gap either. Regulatory impact assessment, including the more legally minded Dutch *wetgevingstoets*,¹⁷ checks compliance burden and administrative feasibility — not whether a court exercising independent judgment will uphold the statute’s reading. Checklists from groups like the National League of Cities or ICMA generally become relevant only after litigation has begun, not at the drafting stage, when corrective revision remains comparatively inexpensive. Standard municipal fiscal forecasts track known budget lines but miss the two costs that matter most after *Loper Bright*: the direct cost of defending an ambiguous ordinance on appeal, and the indirect cost of losing conditioned funds tied to that regulation’s validity.

No existing framework is designed to answer the question at the center of this Article: before a city adopts a regulation, how exposed is it to independent judicial reinterpretation — and what can be changed now to reduce that exposure?

1.4 Contribution and Roadmap

This Article proposes an answer: the Judicial Impact Assessment, or JIA.¹⁸ This Article uses “JIA,” “the Protocol,” and “the framework” interchangeably to refer to the same four-block methodology;

¹⁶ Daniel J. Hay & David R. Carpenter, *The Future of State and Local Agency Deference after Loper Bright*, Sidley Austin Client Alert (June 2025).

¹⁷ See *infra* Part III.1.

¹⁸ As used here, JIA denotes a sub-national regulatory-drafting risk methodology, distinct from the identically named concept in Indian legislative-judicial capacity planning. See Report of the Task Force on Judicial Impact Assessment (India, 2008).

where precision matters, it distinguishes JIA as the overall sequence of decisions from the Municipal Litigation Vulnerability Index (MLVI), the scored instrument that sits inside it. JIA consists of four integrated components. The first scans proposed or existing regulatory text for the kind of ambiguity courts are now most likely to resolve against the city. The second profiles litigation and preemption risk using a scoring tool, the Municipal Litigation Vulnerability Index (MLVI). The third forecasts the fiscal and administrative fallout of an adverse ruling, including effects on conditioned federal funding. The fourth focuses on revising the regulation itself, considering severability clauses, safe harbors, and other drafting techniques that reduce litigation risk while preserving the regulation's underlying policy objectives.

The practical premise of JIA is straightforward: identifying a statutory vulnerability before enactment is, as a general matter, less costly than litigating it after the fact. JIA is designed to give city attorneys a structured way to act on that premise — not a guarantee of litigation avoidance, but a discipline for surfacing the same interpretive weaknesses a reviewing court will eventually look for, before that court does.

The remainder of this Article proceeds as follows. Part II lays out the doctrinal ground *Loper Bright* created, including the return of *Skidmore* and the actual, rather than assumed, degree of divergence across circuits. Part III explains why existing impact-assessment tools, foreign and domestic, miss judicial risk specifically. Part IV builds out the four JIA components in full, including its limits and the objections a careful reader should raise. Part V applies JIA retrospectively to *Sackett* and *San Francisco v. EPA*, then draws a narrow comparative lesson from Ukraine's post-2014 decentralization reform. Part VI sketches an implementation path scaled to cities of different sizes, and closes with an agenda for testing the MLVI empirically. Part VII concludes.

This Article deliberately confines its scope to municipalities, though the same four-block logic could extend to counties, school districts, and other sub-state bodies that draft regulations under similarly borrowed statutory authority. Municipalities are the right place to start not because the problem is unique to them, but because they combine the least institutional capacity with the most direct exposure to the judicial lottery Part II describes — the hardest case, not the only one.

1.5 A Note on Method

This is a doctrinal and conceptual Article, not an empirical one. It does not test JIA against a dataset of adopted ordinances. The MLVI is offered here as a first-generation scoring tool, built from doctrinal analysis and practitioner experience, not statistical calibration.

The two case studies in Part V were chosen because they bracket the range of outcomes cities now face — one adverse, one favorable — not because they represent a random sample of post-*Loper Bright* litigation. Part VI proposes empirical testing of the MLVI against a larger set of municipal regulatory histories as the next stage of this research. This Article is intended to open that inquiry rather than resolve it conclusively.

II. THE POST-LOPER BRIGHT LANDSCAPE: FROM DEFERENCE TO INDEPENDENT INTERPRETATION

2.1 The Scholarly Debate

Loper Bright is barely two years old, and the scholarship it has produced has not yet settled into fixed positions. Thomas Merrill’s warning is the most direct: because the “best reading” of a statute is, in practice, whatever a given panel of judges decides it is, he expects independent interpretation across ninety-four federal districts and thirteen circuits to reproduce, at the judicial level, the instability once blamed on agencies under *Chevron*.¹⁹

¹⁹ Thomas W. Merrill, *The Demise of Deference — And the Rise of Delegation to Interpret?*, 138 Harv. L. Rev. 227 (2024) (Case Comment).

Daniel Hemel's response does not dispute Merrill's mechanism so much as reframe the accounting behind it. Hemel casts the shift in temporal and spatial terms: under *Chevron*, an agency's reading of an ambiguous statute could survive judicial review in every circuit while it held the position, only to flip once a new administration took a different view of the same provision — a regime of uniformity across circuits purchased at the cost of instability over time. *Loper Bright* trades that instability for a different kind. Once a court settles on a statute's best reading, that reading holds until Congress itself amends the statute; what varies now is not the reading over time within a circuit, but the reading across circuits at the same moment, since each circuit is independently tasked with finding its own best answer to the same ambiguous text.²⁰

Hemel compares the two regimes along four dimensions — economic efficiency, regulatory efficacy, learning, and individual choice — and reaches a tentative conclusion that the trade favors *Loper Bright*: spatial disagreement among circuits can eventually be resolved through Supreme Court review, while an agency's own temporal reversals answer to no comparable external check.²¹ He is careful to note the limits of this conclusion; a differently structured statute could in principle produce fewer splits under *Loper Bright* than flips under *Chevron*, and his account is offered as a general tendency rather than a law that holds in every case.²² For a municipality, though, the general tendency is what matters: Hemel's own analysis concedes that where splits do occur, many will persist for years, given the Supreme Court's limited certiorari docket — precisely the interval in which a municipality's ordinance sits exposed.

²⁰ Daniel J. Hemel, *Flips and Splits in Administrative Law*, 74 Duke L.J. 1703 (2025).

²¹ *Id.* (comparing temporal and spatial disuniformity along dimensions of economic efficiency, regulatory efficacy, learning, and individual choice, and concluding that *Loper Bright*'s recalibration toward splits and away from flips is, tentatively, a desirable direction).

²² *Id.* (acknowledging that a differently structured statute could in principle yield fewer splits under *Loper Bright* than flips under *Chevron*, and offering the analysis as a general tendency rather than a universal rule).

Cass Sunstein sits closer to the middle of this exchange. He is willing to treat *Loper Bright* as a genuine constitutional correction — comparing it, in one account, to *Marbury v. Madison*’s assertion that it is the judiciary’s province to say what the law is — yet declines to forecast how much the decision will actually change outcomes in practice, since it preserves *Skidmore* and leaves Congress free to delegate discretion explicitly where it chooses to do so.²³ Jonathan Adler and Ilya Somin, writing separately but in a similar register, caution against overreading the decision in either direction. Adler describes the shift as one from deference to respect rather than a wholesale revolution; Somin’s own assessment is that the ruling is important, but less so than commentary at the time suggested.²⁴

None of this literature was written with a city attorney’s desk in mind, and the empirical work that has followed it has not filled that gap either. The remainder of this Part sets out why that omission matters.

2.2 The Return of *Skidmore*: A Weaker Standard, Unevenly Applied

Loper Bright did not eliminate deference outright. It replaced *Chevron*’s mandatory rule with *Skidmore*’s older, discretionary one, under which an agency’s interpretation earns respect in proportion to its own persuasive force — the thoroughness of its reasoning and its consistency over time — rather than by virtue of coming from an agency at all.²⁵ For a municipality, the distinction has practical consequences. The administrative record behind an ordinance, including its findings and the reasoning documented at the time of adoption, now performs real interpretive work, since that record is precisely what *Skidmore* asks a reviewing court to weigh. Under

²³ Cass R. Sunstein, *Our Marbury: Loper Bright and the Administrative State*, 74 Duke L.J. (forthcoming 2025) (originally circulated as *The Consequences of Loper Bright*, Harvard Public Law Working Paper No. 24-29, July 2024).

²⁴ Jonathan H. Adler, *From “Deference” to “Respect” — The Real Import of Loper Bright*, Volokh Conspiracy (July 3, 2024); Ilya Somin, *The Supreme Court’s Decision Overruling Chevron Is Important — But Less So Than You Might Think*, Volokh Conspiracy (June 28, 2024).

²⁵ *Skidmore v. Swift & Co.*, 323 U.S. 134, 140 (1944); *Loper Bright Enters. v. Raimondo*, 603 U.S. 369, 394–95 (2024).

Chevron, an agency could prevail despite a thin record so long as its interpretation cleared the statute’s ambiguity threshold; *Skidmore* offers no comparable floor.

How much protection *Skidmore* now actually supplies remains an open question, and the disagreement about it is itself worth examining. The Supreme Court cited *Skidmore* approvingly nine times in *Loper Bright*, across the majority opinion and Justice Gorsuch’s concurrence, a pattern one commentator has read as a signal that the Court expected the doctrine to endure.²⁶ Yet the Court did not use the word “*Skidmore*” in a single opinion during the year that followed, including in decisions where its logic has been argued to be operating beneath the surface.^{27, 28} Lower courts have not resolved the question either. The Sixth and Eleventh Circuits, the latter governing Florida, have issued divided opinions on whether *Skidmore* survives *Loper Bright* at all. Most other circuits continue to invoke it, though with declining frequency in fields such as environmental regulation, which generate a substantial share of municipal permitting disputes.²⁹ An independent empirical review of the decision’s first six months reaches a similar conclusion by a different route: lower federal courts are, in that reviewer’s phrase, already diverging on what *Loper Bright* means for *Skidmore* and the related *Auer* doctrine, with no sign yet of convergence.³⁰

²⁶ Abbe R. Gluck, *A Year After Loper Bright: Textualism, Shadow Skidmore, and a New Major Questions Exception*, SCOTUSblog (Oct. 16, 2025).

²⁷ *Kennedy v. Braidwood Management, Inc.*, 606 U.S. 748 (2025) (holding that U.S. Preventive Services Task Force members are inferior officers consistent with the Appointments Clause; not itself a *Skidmore* or statutory-interpretation holding).

²⁸ Gluck, *supra* note 25 (reading *Braidwood*’s discussion of “considered and consistent” agency practice as an indirect, “shadow” application of *Skidmore* logic); see also Ryan P. Mulvey, *No Signs of a Skidmore Revival at the Supreme Court*, Americans for Prosperity Foundation (July 17, 2025) (disputing this reading and arguing the Court’s post-*Loper Bright* references to agency practice are better understood as evidence of original public meaning, not *Skidmore* deference).

²⁹ Gluck, *supra* note 25 (identifying the Sixth and Eleventh Circuits as divided on whether *Skidmore* survives *Loper Bright*, and noting declining invocation of *Skidmore* in environmental and health-care cases relative to immigration and labor cases).

³⁰ Robin Kundis Craig, *The Impact of Loper Bright v. Raimondo: An Empirical Review of the First Six Months*, 109 Minn. L. Rev. 2671, 2673 (2025).

The practical stakes for a municipality are narrower than this doctrinal disagreement might suggest. A well-developed administrative record remains worth building regardless of how the *Skidmore* question is eventually resolved; nothing in *Loper Bright* suggests that a reviewing court should ignore it. What no city attorney can yet promise a client is that the same record will be weighed the same way from one circuit to the next. That uncertainty about how a court will actually use the record a city builds is the condition this Article’s methodology is built to address.

Read together, the disappearance of *Chevron*, the contested revival of *Skidmore*, and the resulting disagreement among circuits amount to more than a change in how courts review regulations after they take effect. They change what a careful regulator needs to do before a regulation is adopted. Section 2.3 turns to how measurable that change already is.

2.3 The “Judicial Lottery” Effect

Independent judicial interpretation, applied by different panels to different statutes, has already produced divergent outcomes as a matter of routine, and the scale of that divergence, though still provisional, can be measured.

Chevron was cited more than 18,000 times over its forty-year life, an average of roughly 450 citations a year.³¹ *Loper Bright* was cited more than 400 times in its first six months alone, a pace suggesting that lower courts are applying it as a matter of course rather than occasionally.³² Robin Craig’s empirical review of that six-month period found that, within her dataset, federal courts invalidated newly promulgated agency rules at a markedly higher rate than they treated other agency actions, which her data show being upheld in a clear majority of cases.³³ That finding rests

³¹ *Id.* at 2673 (reporting that *Chevron* was cited over 18,000 times during its forty-year existence).

³² *Id.* (reporting more than 400 citations to *Loper Bright* within six months of the decision).

³³ *Id.* (reporting that lower federal courts have, within the author’s dataset and coding methodology, invalidated newly promulgated agency rules at a markedly higher rate than they have treated other agency actions, which the same dataset shows being upheld in a clear majority of cases).

on a single six-month sample and Craig’s own coding decisions about what counts as a new rule rather than routine agency action, and it should be read as an early signal from one study rather than a settled baseline. Read with that caveat, it still points toward the asymmetry this Article treats as central: courts appear, within Craig’s dataset, more willing to scrutinize a new regulation than to revisit an old one, which is exactly the position a municipality occupies each time it adopts, rather than merely enforces, an ordinance.

Two pairs of circuit decisions make the same point concretely, in different regulatory settings. The first involves a single agency reviewed by a single circuit specializing in its appeals. Any employer challenging a National Labor Relations Board decision must appeal to a circuit with jurisdiction over the case, but the D.C. Circuit hears the largest share of Board appeals and effectively sets the tone for how much deference the Board receives nationally. Days after *Loper Bright* came down, the D.C. Circuit reaffirmed that Board decisions warrant a “very high degree of deference,” setting one aside only where it departs from established precedent without reasoned justification.³⁴ A month later, the Sixth Circuit confronted the same underlying question — how much independent scrutiny the Board’s interpretation of the same statute should receive after *Loper Bright* — and reached the opposite conclusion, holding it would not defer to the Board’s reading of the NLRA at all, but would exercise independent judgment in deciding whether the Board had acted within its statutory authority.³⁵ Same agency, same statute, one month apart, opposite postures.

The second pair involves different agencies but the same dynamic. The Sixth Circuit set aside the FCC’s 2024 net neutrality order, holding that *Loper Bright* let it resolve, rather than continue to

³⁴ *Hospital de la Concepcion v. NLRB*, 106 F.4th 69 (D.C. Cir. 2024) (reaffirming a “very high degree of deference” to Board decisions days after *Loper Bright*).

³⁵ *Rieth-Riley Constr. Co. v. NLRB*, 114 F.4th 519, 528 (6th Cir. 2024) (declining to defer to the Board’s interpretation of the NLRA and exercising independent judgment instead, one month after *Hospital de la Concepcion*).

defer to, an issue on which the agency had reversed position under successive administrations.³⁶ The Fifth Circuit, reviewing a Department of Labor rule on tipped-employee compensation that had survived judicial review under *Chevron* just years earlier, vacated it once compelled to apply independent judgment to the same ambiguous statutory term — one of the first confirmed instances of an agency action that was lawful under *Chevron* becoming unlawful under *Loper Bright* on materially the same record.³⁷ Neither circuit in either pair was wrong on its own terms. What matters is that four federal appellate panels, working in two unrelated regulatory areas, produced two sets of opposite results applying the same Supreme Court precedent to comparable statutory questions.

This Article calls that pattern the judicial lottery. The point is not that courts are unpredictable in some general sense, but that the same interpretive method, applied in good faith by different panels, now regularly produces different answers to comparable questions, and that a municipality drafting an ordinance today has no reliable way of knowing in advance which panel's reading it will eventually face. A federal agency defending a single nationwide rule encounters this variance once, in whichever circuit a challenger happens to choose. Municipal governments encounter the same variance many times over: each city faces its own single roll of the dice, but the country's roughly 90,000 municipalities collectively face it across all thirteen federal circuits, many times a year.

Part IV returns to this problem directly. JIA's second block gives a municipality a structured basis for estimating, before a regulation is adopted, how its own circuit and its own statutory hook are likely to fare in this still-unsettled landscape.

³⁶ *Ohio Telecom Ass'n v. FCC (In re MCP No. 185)*, 124 F.4th 993, 999 (6th Cir. Jan. 2, 2025).

³⁷ *Restaurant Law Center v. U.S. Dep't of Labor*, No. 24-50186 (5th Cir. 2025) (vacating the Department of Labor's tip-credit rule, which had previously survived judicial review under *Chevron*, upon independent review required by *Loper Bright*).

2.4 The State-Court Dimension

The judicial lottery Part 2.3 describes does not stop at the federal courthouse door. *Chevron* was always a rule about federal statutes; state courts were never bound by it, and *Loper Bright* itself rests on the Administrative Procedure Act, a federal statute with no direct claim on how a state court reads a state law.³⁸ Many state courts nonetheless treat federal administrative-law doctrine as persuasive authority for their own practice, and the result is a second axis of divergence layered on top of the first.

Counts of how many states extend deference to their own agencies vary by source, and the variance is informative in its own right. One practitioner tracker counts at least thirty-four states as extending some level of deference, a broad definition that includes partial or inconsistently applied practice.³⁹ Ballotpedia’s own tracker counts twenty-seven states with what it calls “extant judicial deference doctrines,” a narrower standard requiring a formally established rule rather than scattered instances.⁴⁰ A third survey, focused specifically on courts it describes as employing “robust” deference, counts twenty-five states outright and a twenty-sixth, Utah, applying a qualified version limited to discretionary decision-making.⁴¹ The three counts are not in factual tension with one another. They are answering three progressively narrower questions, and the fact that the same underlying landscape yields thirty-four, twenty-seven, or twenty-five depending on how strictly “deference” is defined is itself a small illustration of the interpretive instability this Article treats as its subject. Whatever the precise count, the direction of change in 2025 is not in

³⁸ *Litigation Minute: A Year After Loper Bright, Part II: States Follow Suit*, K&L Gates (Oct. 13, 2025).

³⁹ *Id.* (counting at least thirty-four states as extending some level of deference to state agencies).

⁴⁰ Ballotpedia, *State Responses to Judicial Deference* (2025) (counting twenty-seven states with what it terms “extant judicial deference doctrines” as of 2025).

⁴¹ Judicial Deference to Agency Expertise in the States, State Court Report (2023, updated 2025) (identifying twenty-five states whose courts “tend to employ robust deference” to agency interpretations, and a twenty-sixth, Utah, applying a narrower, discretion-specific form of deference).

dispute. North Carolina’s Supreme Court, Missouri’s legislature, Louisiana’s governor, Oklahoma’s governor, Texas’s governor, and Kentucky’s legislature, overriding a gubernatorial veto to do so, each eliminated or sharply curtailed deference to state agencies within a single calendar year.⁴² Some states moved by judicial decision, in North Carolina’s case a closely divided one; others by ordinary legislation. Kentucky’s new law states directly that “[a]n administrative body shall not interpret a statute or administrative regulation with the expectation that the interpretation of the administrative body is entitled to deference from a reviewing court.”⁴³

Hawaii’s Supreme Court moved in the opposite direction, declining to follow *Loper Bright*’s reasoning at all and defending its own deferential practice in explicitly comparative terms, contrasting the “experts” federal courts now claim to be with the agency officials Hawaii courts continue to trust with specialized subject-matter knowledge.⁴⁴ A municipality operating under a state enabling statute now faces this same divergence one level down from where Part 2.3 locates it: not only can two federal circuits read an identical federal term differently, but the state court sitting above a given municipality may be moving toward independent review, away from it, or holding steady, largely independent of what any federal circuit does with the same interpretive question. For an ordinance that draws its authority from state, rather than federal, enabling language, this state-court dimension is not a secondary consideration. It is the primary one.

⁴² K&L Gates, *supra* note 37; Ballotpedia, *Four States Have Ended Judicial Deference in 2025* (June 2025); Ballotpedia, *Chevron Deference (Doctrine)* (describing North Carolina Supreme Court’s *Thurman Crofton Savage v. N.C. Dep’t of Transportation* (Aug. 22, 2025) (5–2 decision moving to de novo review); Missouri S.B. 221 (signed July 11, 2025); Louisiana H.B. 99 (signed June 8, 2025); Oklahoma H.B. 2729 (signed May 5, 2025); Texas S.B. 14 (signed Apr. 23, 2025); Kentucky S.B. 84 (veto override Mar. 27, 2025)).

⁴³ Kentucky S.B. 84 (2025).

⁴⁴ K&L Gates, *supra* note 37 (quoting the Hawai’i Supreme Court’s rejection of *Loper Bright*’s reasoning as resting on the premise that Hawaii courts defer “to those agencies with the na’auao (knowledge/wisdom) on particular subject matters to get complex issues right”).

2.5 The Corner Post Problem: No Settled Age for a Regulation

The divergence described so far runs across space — circuit to circuit, state to state. A third dimension runs across time, and it arrived within days of *Loper Bright* itself. On July 1, 2024, three days after deciding *Loper Bright*, the Supreme Court held in *Corner Post, Inc. v. Board of Governors of the Federal Reserve System* that a claim under the Administrative Procedure Act does not accrue, for purposes of the six-year statute of limitations governing suits against the government, until the particular plaintiff bringing it is actually injured by the agency action, regardless of how long ago that action became final.⁴⁵ A regulation adopted decades ago and never successfully challenged is not, on that account, safe. It has simply not yet met the plaintiff who will challenge it.

The Court reached this holding by a 6–3 vote, over a dissent warning that the decision, read together with *Loper Bright*, would unleash what Justice Jackson called a “tsunami of lawsuits against agencies.”⁴⁶ Whatever one makes of that prediction, the mechanism is straightforward and applies with equal force to municipal regulation built on a federal or state statutory hook: a business formed next year, injured for the first time by a decades-old enabling statute or grant condition a municipality’s ordinance depends on, may challenge that authority on a clean slate, with no accrued limitations defense available to shorten the fight.

For a municipality, *Corner Post* removes an assumption that once offered real comfort: that a provision’s age was itself a form of protection, since anyone with standing to challenge it should have done so already. Litigation history under Block 2’s D3, discussed in Part IV, has always

⁴⁵ *Corner Post, Inc. v. Bd. of Governors of the Fed. Rsrv. Sys.*, 603 U.S. 799 (2024) (holding that an APA claim accrues, for purposes of the six-year statute of limitations under 28 U.S.C. §2401(a), when the plaintiff is injured by the challenged agency action, not when the action becomes final).

⁴⁶ *Id.* (Jackson, J., dissenting, joined by Sotomayor and Kagan, JJ.) (warning that *Corner Post*, read together with *Loper Bright*, decided three days earlier, would produce a “tsunami of lawsuits against agencies”).

looked backward, at what circuits and courts have already said about comparable language. *Corner Post* means an ordinance’s clean litigation history to date is not evidence that it will keep that history. It only means no injured plaintiff has yet walked through the door.

2.6 The Major Questions Doctrine: A Layer Above Ordinary Ambiguity

One further doctrine deserves separate treatment, because it operates as something more specific than the ordinary ambiguity Block 1 screens for. The major questions doctrine, given its fullest articulation two years before *Loper Bright* in *West Virginia v. EPA*, holds that where an agency claims regulatory authority over a matter of “vast economic and political significance,” a court will not accept an otherwise plausible reading of an ambiguous statute as sufficient; the agency must instead point to clear, specific congressional authorization for that particular exercise of power.⁴⁷ Justice Gorsuch’s concurrence set out three markers of a major question: an agency asserting power of great political significance, an agency reaching a significant portion of the national economy, or an agency intruding into “an area that is the particular domain of state law.”⁴⁸ That third marker cuts in a direction this Article has not yet had occasion to note. Zoning, land use, and local licensing are, by long tradition, precisely the domain of state and local law the third *West Virginia* factor is protecting. A federal grant condition or agency rule that reaches into local land-use authority is, on the major questions doctrine’s own terms, more vulnerable to challenge for that reason, not less. The doctrine that has narrowed federal agency power generally can, in this specific setting, work in a municipality’s favor against the federal government, even as *Loper Bright* and *Corner Post* work against the municipality’s own exercise of delegated or borrowed

⁴⁷ *West Virginia v. EPA*, 597 U.S. 697 (2022) (holding that, for agency action of “vast economic and political significance,” a court requires clear congressional authorization rather than accepting a merely plausible reading of an ambiguous statute).

⁴⁸ *Id.* (Gorsuch, J., concurring) (identifying three markers of a major question: an agency claiming power of great political significance, an agency reaching a significant portion of the national economy, or an agency intruding into an area that is the particular domain of state law).

authority. A single ordinance can face the major questions doctrine from two directions at once: as a shield against an overreaching federal condition, and as a source of exposure where the ordinance’s own operative language borrows the kind of broad, unspecific delegation that a court applying the doctrine to local authority might treat the same way it treats an overreaching federal agency. Block 1’s ambiguity scan should flag delegation language broad enough to invite this doctrine’s scrutiny in either direction, not simply flag ambiguity as an undifferentiated category. The doctrine reappeared just over a year later, in a case with no connection to environmental law. In *Learning Resources, Inc. v. Trump*, decided February 20, 2026, the Court held that the International Emergency Economic Powers Act does not authorize the President to impose tariffs, applying the major questions doctrine to a claimed delegation of Congress’s core taxing power.⁴⁹ The concurrences split on a question this Article’s own hedging already anticipates. Justice Barrett described the doctrine as “an ordinary application of textualism,” not a freestanding override of the statute’s best reading.⁵⁰ Justice Kagan, joined by Sotomayor and Jackson, agreed with the result but declined to invoke the doctrine at all, concluding that “the ordinary tools of statutory interpretation amply support” the same outcome without it.⁵¹ The disagreement matters because it illustrates that uncertainty can arise not only from ambiguous statutory language, but from disagreement over which interpretive methodology a court should apply to that language in the first place. Block 1 is built to flag exactly that kind of uncertainty, without assuming that one methodological approach has already prevailed.⁵²

⁴⁹ *Learning Resources, Inc. v. Trump*, 607 U.S. ____ (2026) (No. 24-1287) (holding that the International Emergency Economic Powers Act does not authorize the President to impose tariffs, with the major questions doctrine portion of the opinion, Parts II-A-2 and III, joined only by Roberts, C.J., and Gorsuch and Barrett, JJ.).

⁵⁰ *Id.* (Barrett, J., concurring) (describing the major questions doctrine as “an ordinary application of textualism”).

⁵¹ *Id.* (Kagan, J., concurring in part and concurring in the judgment, joined by Sotomayor and Jackson, JJ.) (concluding that “the ordinary tools of statutory interpretation amply support” the result without resort to the major questions doctrine).

⁵² This discussion reflects the state of tariff-related litigation as of July 2026. The Executive Branch’s shift toward Section 122 of the Trade Act of 1974 as an alternative tariff authority, and new challenges filed in the Court of

III LIMITATIONS OF EXISTING REGULATORY IMPACT TOOLS

3.1 RIA and European Models: The Dutch *Wetgevingstoets* and Why a Compliance-Focused Approach Is Insufficient Under American Judicial Activism

Regulatory impact assessment originated in the United States, in the inflation-impact reviews the Carter administration required starting in 1978 and the cost-benefit framework the Reagan administration made mandatory soon after. The idea spread well beyond its origin. By 2000, twenty of twenty-eight OECD countries had adopted some version of RIA, and virtually all OECD members use one today.⁵³ Across these systems, RIA asks roughly the same question: whether a proposed rule is justified, proportionate, and adequately costed against the alternatives. That's a valuable question. It isn't the one this Article is asking.

The Netherlands offers the most developed version of that model, and the clearest illustration of why it doesn't transfer. Since 2023, Dutch legislative drafting has run through the *Beleidskompas*, a whole-of-government framework that succeeded the older *Integraal Afwegingskader*.⁵⁴ Embedded within it is the *wetgevingstoets*, administered by the Ministry of Justice and Security, which reviews roughly five hundred bills and delegated regulations a year for administrative and legal quality, including consistency with constitutional and international-law obligations, before they reach the Council of Ministers.⁵⁵ The Dutch government now publishes an annual *Staat van de Wetgevingskwaliteit* tracking how the process performs.⁵⁶

International Trade in the months following *Learning Resources*, remained unresolved at that time and may affect the practical significance of this discussion by the time of publication.

⁵³ OECD, *Regulatory Policies in OECD Countries: From Interventionism to Regulatory Governance* (2002).

⁵⁴ OECD, *Better Regulation Practices Across the European Union 2025: Netherlands* (2025).

⁵⁵ Ministerie van Justitie en Veiligheid, *Staat van de Wetgevingskwaliteit* (2024); cf. Kamerstuk 31731, nr. 1 (Neth. 2008) (background only; predates the 2023 reform).

⁵⁶ Ministerie van Justitie en Veiligheid, *supra* note 53.

This is a strong model, built for a different constitutional structure. The mismatch matters here, not as a curiosity, but as the point itself. The *wetgevingstoets* operates inside a system of parliamentary supremacy, where courts don't independently reinterpret primary legislation against the legislature's evident intent, and where the reviewing body sits inside the same executive that drafts the bill. Dutch courts don't generate the judicial lottery Part II describes. American municipalities face substantially different institutional conditions. Under *Chevron*, this mismatch mattered less, since much of the ambiguity risk a *wetgevingstoets*-style test would have screened for was already absorbed by mandatory agency deference before a court ever saw it. *Loper Bright* removed that buffer. A drafting-quality test built for a system with no judicial lottery was never designed to price one. It therefore cannot answer the question JIA is designed to answer.

3.2 A Federalist Common-Law Alternative: Canada's *Vavilov* Framework

Canada shows the same diagnosis reaching the opposite prescription. In December 2019, the Supreme Court of Canada found that *Dunsmuir v. New Brunswick*, the framework then governing judicial review, had failed to deliver the simplicity and predictability it promised.⁵⁷ A seven-member majority in *Canada (Minister of Citizenship and Immigration) v. Vavilov* didn't abandon deference in response. It sharpened it. Reasonableness became the presumptive standard across administrative decision-making, with correctness review confined to a statutory right of appeal or one of three narrow categories: constitutional questions, general questions of law central to the legal system, and jurisdictional boundaries between administrative bodies.⁵⁸ The Court later extended this framework to subordinate legislation, the closest Canadian analogue to a municipal

⁵⁷ *Canada (Minister of Citizenship and Immigration) v. Vavilov*, 2019 SCC 65 (quoting *Dunsmuir v. New Brunswick*, 2008 SCC 9).

⁵⁸ *Vavilov*, 2019 SCC 65; *Bell Canada v. Canada (Attorney General)*, 2019 SCC 66 (companion case decided the same day).

bylaw.⁵⁹ The parallel to *Loper Bright* holds only within a narrower band than it first appears. *Chevron* and *Loper Bright* governed one specific question: how much weight a court owes an agency's reading of ambiguous statutory text. *Vavilov*'s reasonableness standard covers more ground, and statutory questions simply fall within that wider sweep. Where the two doctrines do overlap, they point in opposite directions. The U.S. Supreme Court eliminated the deferential default for statutory questions; Canada kept it and confined independent review to three closed categories. Canada's fix doesn't solve the American municipal problem, and not because Canada carries less uncertainty. *Vavilov*'s categories are settled by one national Supreme Court sitting atop a small number of intermediate courts, which gives them room to converge nationally over time.⁶⁰ American municipalities have no equivalent single interpreter to settle anything. Canada demonstrates that abandoning deference is not the inevitable response to concerns about unpredictability; preserving deference is itself a viable institutional choice.

3.3 A Common-Law System That Never Adopted Deference: Australia

Australia offers a third data point, further along the same spectrum. Its High Court has never adopted anything like *Chevron* deference, treating statutory interpretation as an exclusively judicial function under a constitutionally entrenched separation of powers, and confining deference to *Wednesbury*-style review of discretionary and factual findings rather than legal ones.⁶¹ Australian courts have lived with something structurally close to *Loper Bright*'s independent-interpretation default for decades, not two years. If that experience holds a lesson for American

⁵⁹ See Supreme Court of Canada Holds That the *Vavilov* Framework Applies to the Review of Subordinate Legislation, Dentons Commercial Litigation Blog (2024).

⁶⁰ Dentons Canada Regulatory Review, *Judicial Review Post-Vavilov: Are Courts "Correcting" Administrative Decisions?* (2022) (surveying post-*Vavilov* decisions).

⁶¹ Janina Boughey, *A Perspective from a Jurisdiction Without a Doctrine of Deference: Australia*, Australian Public Law (2023).

municipalities, it's a sobering one rather than a reassuring one. Even a mature common-law system with a century of accumulated precedent hasn't eliminated the unpredictability *Loper Bright* now introduces to American statutory interpretation. The predictability Australian courts enjoy does not come from the absence of *Chevron* as such. It comes from a century of accumulated practice giving content to what independent interpretation means in any given context. American municipalities face the same interpretive model without the benefit of that accumulated practice — closer to what this Article proposes JIA start doing now than evidence that the instability resolves on its own with time.

3.4 A Civil-Law Federation Without Deference: Germany

Germany adds a further, structurally distinct data point: a civil-law federal system with a written constitution, where courts have likewise never developed anything resembling *Chevron* deference. German administrative courts review legal questions in full as a matter of course, reserving deference only for narrow categories of specialized technical judgment or genuine policy discretion.⁶² Neither covers statutory interpretation itself.

Germany adds something Australia doesn't: a second variable this Article has treated as separate from judicial deference but structurally related to it. Article 28 of the Basic Law guarantees municipal self-government directly, and German municipalities may bring a constitutional complaint to the Federal Constitutional Court alleging that a federal or state law has infringed that guarantee.⁶³ Most American municipalities, operating under the Dillon's Rule regimes discussed

⁶² Jan S. Oster, *The Scope of Judicial Review in the German and U.S. Administrative Legal System*, German Law Journal (2010).

⁶³ Grundgesetz [GG] [Basic Law], art. 28, art. 93(1)(4b) (Ger.); Bundesverfassungsgerichtsgesetz [BVerfGG] § 91 (Ger.) (“Gemeinden und Gemeindeverbände können die Verfassungsbeschwerde mit der Behauptung erheben, daß ein Gesetz des Bundes oder des Landes die Vorschrift des Artikels 28 des Grundgesetzes verletzt.”) (providing that municipalities and municipal associations may bring a constitutional complaint alleging that a federal or state law violates Article 28's self-government guarantee, subject to subsidiarity where equivalent protection exists under state constitutional law).

in Part I, have markedly weaker constitutional protection for local self-government than German *Gemeinden* do. Home-rule states offer some floor, but nothing resembling a directly enforceable constitutional guarantee reachable through a complaint to the country's highest court.

Even so, Germany's experience doesn't reproduce the judicial lottery Part II describes. The Federal Constitutional Court is one court, not thirteen independent circuits, so full judicial review of legal questions there doesn't fragment into circuit-by-circuit divergence. Full review and unpredictable review aren't the same thing, and Germany shows the first can exist without the second. The American problem therefore lies not in independent judicial review itself, but in the institutional fragmentation within which that review occurs. These two exposures — a thinner deference cushion and weaker constitutional protection for municipal authority — are independent problems that happen, in the American case, to have arrived together.

3.5 Ex Ante Screening Where the American Constitution Forbids It: France

France adds a final, sharply different data point: a system that builds ex ante judicial screening directly into the legislative process, something the American constitutional structure forecloses outright. Before a government bill reaches Parliament, the Conseil d'État — sitting in an advisory rather than adjudicative capacity — must review it, including its constitutional soundness.⁶⁴ After Parliament votes but before the President promulgates the law, the Conseil constitutionnel conducts a separate review of its own, mandatory for some categories of legislation and discretionary for others, on a one-month clock.⁶⁵ A law found unconstitutional cannot be

⁶⁴ *Conseil constitutionnel, La mise en place de la révision constitutionnelle en France et la question préjudicielle de constitutionnalité* (describing the Conseil d'État's mandatory prior advisory review of government bills).

⁶⁵ *Conseil constitutionnel, Présentation générale*; *Sénat, Le contrôle de constitutionnalité des lois* (describing *contrôle a priori* under Article 61 of the Constitution).

promulgated at all. Some of the judicial risk this Article is concerned with is filtered out before the statute ever takes effect.

No comparable option exists for an American municipality, and not for lack of institutional imagination. The case-or-controversy requirement of Article III has long forbidden American courts from issuing advisory opinions on unenacted or unchallenged law.⁶⁶ The American system cannot import the French solution; the obstacle is constitutional, not institutional. A municipality can't ask a court what an ordinance means before adopting it, no matter how much it might want to. That constitutional foreclosure is precisely why an ex-ante instrument like JIA has to be built by the regulated party itself, in the absence of any judicial body positioned to do the screening in advance.

3.6 Why These Five Systems

The five systems compared above weren't chosen to survey comparative administrative law broadly. Each isolates one variable this Article treats as analytically distinct from the others. The table below summarizes what each system contributes. Other systems would likely offer further variations on these same themes, and this Article doesn't claim the five selected are the only ones that could. They were chosen for the sharpness of the contrast each supplies to the American case, not for geographic or doctrinal completeness.

System	Legal family		Deference posture		Isolated variable	
Netherlands	Civil	law,	Not	applicable	—	Whether the underlying predictability
	unitary		parliamentary	supremacy		problem exists at all

⁶⁶ *Muskrat v. United States*, 219 U.S. 346 (1911) (grounding the prohibition on advisory opinions in Article III's case-or-controversy requirement).

System	Legal family	Deference posture	Isolated variable
Canada	Common law, federal	Recalibrated toward more deference (<i>Vavilov</i>)	Whether a court can choose deference over independent review after diagnosing instability
Australia	Common law, federal	Never adopted deference	What independent interpretation looks like over decades, not two years
Germany	Civil law, federal	Never adopted deference	Whether full review necessarily produces a judicial lottery, or the two can come apart
France	Civil law, unitary	Ex ante screening; deference question does not arise	Whether ex ante judicial screening is structurally possible at all

Together, these comparisons suggest that JIA addresses a distinctly American institutional problem rather than a generic feature of administrative law.

Nothing here suggests the Dutch, Canadian, Australian, German, or French instruments are deficient. Each does its own job well. The point is narrower and, this Article thinks, more useful: legality, proportionality, economic cost, and administrative feasibility are the dimensions regulatory-quality review has traditionally measured, at home and abroad, and every system canvassed above measures some combination of them competently. None was built to measure whether a regulation's own text can survive independent judicial reinterpretation by a body the drafter does not control. That is a different evaluative question, not a deeper version of the ones these tools already answer, and Part IV's methodology is offered as a first attempt to answer it in a structured way.

3.7 Beyond Fiscal Impact: Why Conventional Budgetary Forecasts Miss Litigation Costs and Grant Risk

Domestic tools close the gap no better, for different reasons. U.S. municipal fiscal forecasting is a mature practice, built around known and recurring budget lines: personnel costs, debt service, capital projects, pension obligations. Standard public-finance guidance treats the kind of cost this Article is concerned with only at the most general level. The leading recommendation on municipal reserves instructs governments to hold sufficient unrestricted fund balance to absorb “unpredicted one-time expenditures,” without distinguishing a specific, statute-linked risk, such as an adverse circuit ruling on an ordinance’s own legal basis, from an ordinary revenue shortfall.⁶⁷ That framework was built for volatility in tax receipts and service demand, not for the contingency this Article addresses.

Two costs fall into that gap. The first is the direct cost of litigation itself. *City and County of San Francisco v. EPA* shows the scale this can reach even for a well-resourced city. The underlying permit dispute ran roughly six years, from the contested 2019 permit terms through denial by the Environmental Appeals Board, an unsuccessful petition to the Ninth Circuit, and finally Supreme Court review. The city’s own filings estimated that losing the case would have required at least \$10 billion in capital expenditures, pushing ratepayer utility bills toward \$9,000 a year — the city’s own advocacy estimate rather than an independently audited number, but illustrative of a scale no ordinary budget line anticipates.⁶⁸

⁶⁷ Government Finance Officers Association, *Fund Balance Guidelines for the General Fund* (recommending, per Nat’l Advisory Council on State and Local Budgeting, Recommended Practice 4.1, that governments maintain reserves against “unpredicted one-time expenditures”).

⁶⁸ *City and County of San Francisco v. EPA*, 604 U.S. 334 (2025); San Francisco Public Utilities Commission, *San Francisco Files Final Brief in SCOTUS Case to Protect Utility Ratepayers from Massive Bill Increases* (2024).

The second cost is risk to conditioned federal funding. Roughly a quarter of all state and local transportation and water infrastructure spending nationally is paid for through federal grants.⁶⁹ When the underlying federal authority for a grant condition is itself narrowed by post-*Loper Bright* litigation, a municipality that built its budget around that funding stream inherits the disruption without having been a party to the case that caused it.

A related but distinct risk runs through delegated regulatory authority rather than a grant condition as such. Florida's experience with wetlands permitting illustrates the point. When a federal appellate court held that EPA and the Fish and Wildlife Service had unlawfully approved the state's delegated Section 404 permitting program under the Endangered Species Act, permitting authority reverted to federal agencies.⁷⁰ Every local government that had planned development-review timelines around state-level processing absorbed the resulting delay, without having litigated the underlying question itself. The mechanism differs; the consequence doesn't.

Existing practitioner tools don't price any of these costs, because they weren't built to. The checklists and model ordinances circulated by organizations like the National League of Cities, and by similar practitioner associations, are consulted, as a matter of design, once litigation is filed or threatened.⁷¹ They're response protocols, meant to manage a dispute that has already begun, not forecasting instruments meant to price one that hasn't yet occurred.

No existing framework, domestic or comparative, systematically asks the question this Article treats as central. Before a specific ordinance is adopted, how exposed is its specific statutory language to independent judicial reinterpretation, and what would that exposure cost if it

⁶⁹ The Pew Charitable Trusts, *State and Local Governments Face Persistent Infrastructure Investment Challenges* (2023).

⁷⁰ See *supra* Part I, note 7 (Earthjustice, *Court Victory for Florida Wetlands Protection* (2026)).

⁷¹ See *supra* Part I, note 13 (citing National League of Cities model-practice materials).

materialized? Part IV develops a methodology built to ask exactly that question, at the only point in the process, before adoption, when the answer can still change the outcome.

IV. THE JUDICIAL IMPACT ASSESSMENT FRAMEWORK: A FOUR-BLOCK APPROACH

4.0 Design Logic

Part III named the gap the five systems compared there don't fill: none is built to measure whether a regulation's own text can survive independent judicial reinterpretation. This Article argues that judicial resilience should be understood as a distinct dimension of regulatory quality, separate from legality, proportionality, economic efficiency, and administrative feasibility — the dimensions existing review instruments, domestic or foreign, were built to assess. JIA is offered as the first structured methodology for operationalizing that dimension, not as a claim that it is the only evaluative criterion still missing from regulatory review.

JIA rests on four decisions a municipality makes, in sequence, before adopting a regulation. Where does the statute's language admit more than one reading — a question of interpretation. How is a court in this circuit likely to treat that reading — a question of exposure. What would an adverse resolution cost — a question of consequence. And what should be rewritten in response — a question of redesign. Each decision depends on the one before it, which is why the sequence in Figure 1 runs in one direction, with exits at three points: a regulation that is not material, one where Block 1 finds no genuine ambiguity, and one Block 2 scores as low risk.




Figure 1. JIA decision sequence, with exits at each of three points.

Materiality is doing real work here. JIA is meant for regulations whose language is borrowed from a state statute or a federal grant condition of unsettled validity, that carry meaningful fiscal exposure, or that touch a subject — zoning, environmental permitting, utilities, licensing — where post-*Loper Bright* circuit divergence, documented in Part II, is already live. The office expected to run this is the municipal legal function itself: in-house counsel where a municipality has it, retained outside counsel where it doesn't, applied before adoption rather than after a challenge has already been filed. Part VI takes up how this work gets divided as municipal size and staffing change.

JIA is a framework, not a measurement. It is a sequence of four decisions; only the second of them produces a number. The Municipal Litigation Vulnerability Index earns the word “index” honestly — a defined scale, scored dimensions, a composite, threshold tiers — and that is what separates it from the framework around it. JIA organizes decisions. MLVI measures one.

It measures for screening, not for prediction. A score sorts cases for further attention under rules the municipality can see; it does not state a fact about the world with a margin of error attached. A score of sixty-five is not a sixty-five percent chance of losing in court. It means the underlying dimensions place the regulation in a tier that calls for a certain depth of review.

Five dimensions is not an arbitrary number. Every post-*Loper Bright* risk this Article is concerned with comes down to one of two questions: how likely is an unfavorable reading, and how well can the municipality absorb it if one comes. Four dimensions — statutory language, administrative record, circuit history, state preemption — answer the first question. The fifth, fiscal resilience capacity, answers the second. This Article treats those two questions as sufficient to describe the risk it is trying to capture, at least as a starting hypothesis, which is why the model stops at five rather than admitting a sixth, a seventh, or more.

None of this is a JIA invention. Treating risk as likelihood joined to consequence is standard enterprise risk-management practice.⁷² Nor is combining several dimensions into one composite score new outside law: Virginia Apgar’s 1953 newborn assessment score sums five independently rated components — heart rate, respiratory effort, muscle tone, reflex irritability, and color — into a single number still used to triage which infants need immediate care.⁷³ Composite scoring of

⁷² Committee of Sponsoring Organizations of the Treadway Commission, *Enterprise Risk Management—Integrating with Strategy and Performance* (2017) (defining risk assessment as the joint evaluation of likelihood and impact).

⁷³ Virginia Apgar, *A Proposal for a New Method of Evaluation of the Newborn Infant*, 32 *Curr. Res. Anesth. & Analg.* 260 (1953) (introducing a five-component composite score — heart rate, respiratory effort, muscle tone, reflex irritability, and color, each rated 0 to 2 — still in standard clinical use).

ambiguous or discretionary judgment has legal precedent too: the Federal Sentencing Guidelines fold an offense-level score and a criminal-history category into a single number that sets the sentencing range, on the premise that disclosed, structured scoring beats undisclosed discretion.⁷⁴ That precedent carries a well-documented failure mode along with it — critics have argued for decades that its arithmetic can flatten the individualized judgment sentencing is supposed to involve — and MLVI is offered with that failure mode already in view, as a first attempt rather than a settled instrument.⁷⁵

The point ranges marking low, medium, and high risk are illustrative, not derived from data, and they leave some deliberate slack near the boundaries. Their job is administrative triage, not statistical classification, and calibrating them against real outcomes is exactly the empirical work Part VI proposes next. Because every later judgment presupposes an identified ambiguity, textual analysis necessarily comes first: a court cannot be asked how it would treat a reading that competent lawyers do not actually dispute.

4.0.1 Why Four Blocks Are Sufficient

Each block answers a question the one before it cannot. Block 1 asks whether a provision’s language admits more than one reading — a question about the text itself, answerable without reference to any court. Block 2 asks how a specific circuit is likely to treat that reading — a question about institutional context a linguistic reading alone cannot answer. Block 3 asks what an adverse outcome would cost — a fact about the municipality’s own budget, not about the statute. Block 4 asks what to do about it — a drafting question that presupposes answers to the first three.

⁷⁴ U.S. Sentencing Guidelines Manual §5A (U.S. Sentencing Comm’n 2024) (combining an offense-level score and a criminal-history category into a composite sentencing-range determination).

⁷⁵ See, e.g., Kate Stith & José A. Cabranes, *Fear of Judging: Sentencing Guidelines in the Federal Courts* (1998); cf. *United States v. Booker*, 543 U.S. 220 (2005) (rendering the Guidelines advisory).

The blocks are therefore sequential rather than cumulative: each produces information the following block requires but cannot generate for itself. Four is not a preference. It is the number of categorically distinct questions this Article has identified in the post-*Loper Bright* problem.

4.0.2 Why an Index Rather Than Alternative Architectures

Two alternatives are especially relevant. A binary checklist — flagged or not flagged, on each dimension — cannot rank. A legal department with forty ordinances scored medium risk still needs to know which one to redraft first, and a checklist returning five yes-or-no answers gives no basis for that ordering. A composite score does, by compressing five judgments into one number sortable across an entire docket. A branching decision tree fares no better here: it works where legal consequences follow from clearly defined conditions, but independent judicial interpretation after *Loper Bright* rarely proceeds through binary branches, since courts routinely weigh several competing considerations at once — a poor fit for the problem JIA addresses, whatever value a decision tree has for sequencing the four blocks themselves, as Figure 1 does at a different level of the analysis.

A machine-learning model trained to predict litigation outcomes was also considered, and rejected for a different reason. Too few post-*Loper Bright* outcomes yet exist to train one responsibly, and a probability generated by an opaque model would produce exactly the kind of undocumented conclusion *Skidmore* does not reward, however accurate it turned out to be. A rubric a second reviewer can trace against the state’s own statutory text serves the reasoning *Skidmore* asks a record to show. A black-box score, however well it performed statistically, would not.

4.0.3 Residual Uncertainty and the Limits of Prediction

MLVI does not claim to capture every source of unpredictability a municipality faces, and some of what remains outside it is worth naming rather than leaving implicit. A change in the

composition of the reviewing circuit panel, an amendment to the underlying state or federal statute after adoption, a factual record that develops differently in litigation than counsel anticipated at drafting, and the litigation strategy an eventual plaintiff chooses to pursue are all variables no five-dimension score fixed at the moment of adoption can price. This is not a defect particular to MLVI; it is a limit on what any ex-ante instrument can do. A municipality can follow every step of JIA and still lose, just as it can prevail despite a score that flagged substantial vulnerability — neither outcome, by itself, says whether the framework did its job.

4.1 Block 1: Statutory Ambiguity and Interpretation Scan

Block 1 puts one question to every operative provision in a proposed or existing regulation. Using the tools *Loper Bright* now requires courts to use — ordinary public meaning, statutory structure, the whole-act and whole-code presumptions, canons like *noscitur a sociis* and *ejusdem generis*, a representative rather than exhaustive set of what *Loper Bright* now puts in play — would competent lawyers, arguing in good faith, land on different readings? A provision drafted around borrowed language, rather than the municipality’s own, inherits whatever ambiguity already sits in that borrowed text, which is exactly why it deserves priority attention here. Where no real disagreement is plausible, the regulation exits at this stage. Where it is plausible, Block 1 hands Block 2 an inventory: the specific terms in question and the readings each one supports.

4.2 Block 2: Litigation Risk Profile and State-Level Preemption

Block 2 scores each flagged regulation across five dimensions, twenty points apiece, for a composite MLVI score out of one hundred.

D1 — statutory language quality (0–20). How much of the operative text is ambiguous, overbroad, or tied to agency guidance that independent judicial review could now unsettle. Broad, unspecific delegation language scores high here for a second reason beyond ordinary ambiguity:

it is the kind of language the major questions doctrine, discussed in Part 2.6, treats with particular suspicion, in either direction the doctrine happens to run.

D2 — administrative record completeness (0–20). How developed the regulation’s evidentiary foundation is — not its volume, but its thoroughness, internal consistency, and the soundness of its reasoning, the same qualities *Skidmore* asks a reviewing court to weigh.

D3 — litigation history and circuit exposure (0–20). How the municipality’s own circuit, and any circuit whose reasoning it tends to follow, has already treated comparable language. A clean litigation history lowers this score, but *Corner Post*, discussed in Part 2.5, means it lowers it less than it once did: an untested provision’s silence may reflect the absence of an injured plaintiff so far, not settled law.

D4 — state preemption risk (0–20). This dimension splits into three sub-components, scored separately and summed, rather than judged as one impression. Enabling-statute clarity (0–8) asks how explicitly the state statute authorizes the specific regulatory choice — a municipality regulating under a statute that names the subject matter directly scores low here; one reaching the same subject through a general “health, safety, and welfare” grant scores high. Greater weight is assigned to enabling-statute clarity than to the other two sub-components because the absence of delegated authority generally ends the inquiry before broader home-rule or preemption questions ever arise. Home-rule scope (0–6) asks how far the municipality’s home-rule authority extends in this policy area under the state’s own constitutional or statutory framework, independent of the specific enabling statute — broad in a strong home-rule state, thin or absent in a Dillon’s Rule one. Preemption posture (0–6) asks how actively the state legislature has moved, generally, to preempt local regulation in this subject area — a recent override of local ordinances on the same or an adjacent subject signals live risk regardless of what the enabling statute or home-rule doctrine

independently provide. Splitting D4 this way doesn't raise its ceiling. It replaces one impressionistic judgment with three narrower ones a second reviewer can check independently against the state's own texts. The 8/6/6 point allocation is itself illustrative rather than empirically derived, the same discipline this Article applies to the risk tiers generally, and calibrating the relative weight of these three sub-components against real outcomes is further work for the empirical agenda Part VI proposes.

D5 — fiscal resilience capacity (0–20). How much budgetary cushion the municipality holds against an adverse outcome, the consequence side of the ledger described in 4.0.

The five scores sum to a composite that sorts into three tiers: low risk (0–39), where the regulation proceeds with a confirmed severability clause; medium risk (40–74), where the highest-scoring dimensions get redrafted; high risk (75–100), where advancement halts pending full redesign under Block 4. The precise number, not just the tier, still matters for one reason: a legal department screening a full annual portfolio of ordinances needs to know not only which are medium risk, but which medium-risk ordinances to redraft first, and a composite of 74 sitting at the top of that tier deserves attention before one sitting at 41.

Two examples show the instrument doing different work — one worked through completely, one sketched to show a different exposure profile. A coastal municipality's post-*Sackett* wetlands ordinance scores 16 of 20 on D1: its jurisdictional trigger still tracks pre-*Sackett* WOTUS terminology the agency itself abandoned, and the term carries a fractured judicial history running back to *Rapanos*. It scores 10 of 20 on D2: an environmental impact statement exists but predates *Sackett* and hasn't been supplemented to reflect the narrower test the case adopted. It scores 15 of 20 on D3, given an active Eleventh Circuit litigation climate for exactly this kind of provision. It scores 6 of 20 on D4 — 3 of 8 on enabling-statute clarity, because the state statute authorizes

coastal development regulation generally rather than wetlands regulation specifically; 2 of 6 on home-rule scope, reflecting a moderately strong home-rule position; 1 of 6 on preemption posture, because the state hasn't recently moved to override local wetlands ordinances. And it scores 10 of 20 on D5, an adequate but unexceptional reserve fund. The composite is $16 + 10 + 15 + 6 + 10 = 57$, squarely medium risk, and below the sixteen-point mark on every individual dimension, so the composite governs the classification rather than any single flagged score.

Dimension	Score	Basis
D1 — statutory language quality	16 / 20	Pre- <i>Sackett</i> WOTUS terminology; fractured history since <i>Rapanos</i>
D2 — administrative record	10 / 20	EIS exists but predates <i>Sackett</i> 's narrower test completeness
D3 — litigation history and circuit exposure	15 / 20	Active Eleventh Circuit litigation climate
D4 — state preemption risk	6 / 20	3/8 enabling-statute clarity + 2/6 home-rule scope + 1/6 preemption posture
D5 — fiscal resilience capacity	10 / 20	Adequate but unexceptional reserve fund
Composite	57 / 100	Medium risk (40–74)

Medium risk sends this ordinance to Block 4's targeted redrafting rather than the full-scale redesign reserved for high-risk provisions, though a cautious legal department might run Block 3's cost forecast anyway, given the scale *City and County of San Francisco v. EPA* shows for a structurally similar jurisdictional dispute.

A municipality running a Title X-funded program looks different, and shows why D1 measures the provision, not the drafter. Its language incorporates a federal grant condition whose own validity is contested elsewhere, discussed in Part III, and it scores high on D1 no matter how carefully the municipality wrote its own text, because the exposure was never local to begin with. A third profile shows the same architecture working outside environmental law entirely. A municipality regulating short-term rental licensing under a general “health, safety, and welfare” home-rule grant may score only moderate on D1, since “short-term rental” is not itself a term burdened by unsettled federal meaning the way WOTUS or a Title X condition is. D4 can still score high, because several state legislatures have recently moved to preempt local short-term rental regulation altogether, regardless of how carefully the ordinance’s own text is drafted. The risk here runs almost entirely through D4 rather than D1 — a different profile from either environmental example, and evidence that MLVI’s architecture travels across subject areas rather than being built for permitting disputes specifically.

4.2.1 False Positives and False Negatives

A screening instrument earns its keep by being honest about how it fails, not by pretending it doesn’t. MLVI produces two kinds of error a municipality should expect, not merely tolerate. A false positive occurs when a cautious attorney scores an ordinance high risk, the municipality invests in redesign, and a court later upholds the original language anyway — resources spent on a redesign the litigation never demanded. A false negative is the more dangerous failure: a score that lands low or medium because the record looked settled at the time, followed by a circuit reading the same term in a way the score never anticipated. Robin Craig’s dataset, discussed in Part II, illustrates one version of this risk rather than explaining it fully: courts in her sample invalidated new agency rules at a materially higher rate than they revisited old ones, meaning the

record available before adoption is often thinner than the one a court eventually builds after a challenge is filed. The deeper cause runs wider than any single dataset. It is the same set of variables named in 4.0.3 — a changed panel, an amended statute, a factual record that develops unpredictably, a plaintiff’s litigation strategy — arriving after the score was fixed rather than before.

4.2.2 Professional Judgment and Reproducibility

Two attorneys scoring the same ordinance from the same record will not always land on the same composite, particularly on D1 and D3, where the judgment calls are closer to interpretive argument than to counting. This Article does not claim MLVI eliminates that variance. It claims only that a disclosed five-dimension rubric narrows the variance further than an undocumented gut check does. Reproducibility is therefore treated as an empirical property of the instrument to be measured, not an assumption about professional agreement to be asserted — precisely the test the independent-reviewer study proposed in Part VI is designed to run.

4.3 Block 3: Fiscal and Administrative Impact Forecast

D5 is a snapshot, backward-looking and general. Block 3 looks forward at one specific provision, whatever Block 2 flagged as medium or high risk, and asks what it would actually cost if litigated to a loss, using the cost categories Part III laid out. Where a municipality has the staff for it, JIA splits this work rather than piling it all on legal counsel: an attorney scores D5 along with the rest of the MLVI, and the finance or budget office handles Block 3. Smaller municipalities without a separate budget analyst usually run both through the same legal review, and Part VI addresses that scaling directly.

Stage	Institutional owner
-------	---------------------

Blocks 1–2	Municipal legal function
------------	--------------------------

Stage	Institutional owner
Block 3	Finance or budget office, where one exists separately
Block 4	Municipal legal function

A municipality can hold a healthy D5 score and still face a Block 3 forecast large enough to justify redesigning one provision. A municipality with a thin D5 has more reason to treat even a modest Block 3 number as urgent. Keeping the two separate guards against a specific distortion a merged score would invite. A well-resourced city's deep D5 cushion could mask what a single bad provision actually costs, quietly encouraging it to adopt poorly drafted regulations because the reserve fund can absorb the hit regardless. A fiscally thin city, meanwhile, could be talked out of a well-drafted, genuinely low-risk ordinance by a uniformly weak D5 score that has nothing to do with that ordinance in particular. Block 3 asks for no actuarial model, only a documented line from each flagged provision to the budget line or funding stream it would actually hit if things went wrong.

4.4 Block 4: Regulatory Redesign and the Doctrine of Severability

Block 4 takes each high-priority provision and asks what specific change in drafting would lower its exposure without giving up its purpose. Each dimension points toward its own response.

Flagged dimension	Drafting response
D1 — statutory language quality	Safe harbor: a narrower, textually secure compliance path running alongside the broader one
D2 — administrative record completeness	Supplementary findings and hearing record, added before adoption
D3 — litigation history and circuit exposure	Alternative wording that tracks language the circuit has already upheld

Flagged dimension	Drafting response
D4 — state preemption risk	Express state-law authorization, or narrowing to authority no one disputes
D5 — fiscal resilience capacity	No drafting response — it informs which redesigns are urgent, not what they should say

Two techniques cut across dimensions rather than answering to just one. Severability handles provisions Block 1 has already flagged as internally divisible. American courts have long presumed that an unauthorized part of an enactment can be cut away while the rest stands, asking only whether the enacting body would have adopted the valid remainder on its own.⁷⁶ *Barr v. American Association of Political Consultants* reaffirmed that presumption and put weight on the clause’s own text; *Murphy v. NCAA* shows where it runs out, holding provisions inseverable because the line between what was authorized and what wasn’t proved too blurred to draw.⁷⁷ A severability clause is only as good as the drafting behind it, working only if the ordinance’s provisions are actually built to stand alone, which is precisely what Block 1’s scan is meant to check. Layered drafting is the other half of this: pulling a vulnerable provision into its own section rather than burying it inside lower-risk requirements. It isn’t just tidier drafting. Under the kind of independent scrutiny *Loper Bright* now invites, it is the structural condition severability doctrine needs to have anything to work with. In some settings, a sunset clause or a mandatory post-adoption review date can also reduce long-term exposure where the governing doctrine remains genuinely unsettled, giving the municipality a built-in point to revisit a provision before an adverse ruling forces the question.

⁷⁶ See, e.g., *Murphy v. Nat’l Collegiate Athletic Ass’n*, 584 U.S. 453 (2018).

⁷⁷ *Barr v. American Ass’n of Political Consultants, Inc.*, 591 U.S. 610 (2020) (plurality opinion); *Murphy*, 584 U.S. at 482–83.

Block	Question	Output	Decision
1	Is there genuine ambiguity?	Flagged provisions	Exit, or continue to Block 2
2	How exposed is this circuit/statute combination?	MLVI composite and tier	Exit (low risk), or continue to Block 3
3	What would an adverse ruling cost?	Fiscal forecast	Proceeds to redesign, scaled to the forecast
4	What should change in the drafting?	Revised provision, severability clause, or a documented decision not to proceed	Adoption

A concrete example shows what this redesign looks like in practice, rather than only in the abstract. A coastal zone provision drafted before this kind of scan might read: “Development of the coastal zone is permitted provided that there is minimal negative impact on wetlands.” The phrase does real work for the drafter, but none of it survives independent judicial construction — “minimal negative impact” supplies no standard a court can apply without importing its own judgment, and “wetlands” invites exactly the WOTUS-dependent uncertainty Part V’s discussion of *Sackett* traces in detail. A redrafted version, incorporating Block 1’s diagnosis and Block 4’s response, might instead read: “Development is permitted exclusively on land parcels that have no hydrological connection to waters of the United States as defined in *Sackett v. EPA*, 598 U.S. 651 (2023). In the event that any part of this ordinance is declared invalid by a court of competent jurisdiction, the remaining provisions shall remain in full force and effect.” The revised text anchors the operative standard in a binding, already-decided case rather than an agency’s own

gloss on a statute, and adds the severability clause that gives a court something to preserve if it invalidates any single provision. Neither change requires guessing how a future court will rule. Both changes simply stop asking a court to do interpretive work the ordinance's own drafter could have done first.

Sometimes Block 4's honest answer is that a regulation's purpose can't be pursued through its current statutory hook without accepting the risk Block 2 already flagged. Reaching that conclusion before adoption costs far less than reaching the same conclusion after a court has already ruled. That conclusion is itself a successful outcome of the Protocol, not a failure of it.

This raises a fair objection rather than settling it: does a high MLVI score simply push a municipality toward abandoning contested policy areas altogether, rather than regulating them more carefully? The concern is worth taking seriously rather than waving off, and the honest answer has two parts. First, Block 4's ordinary output is redesign, not abandonment — a narrower safe harbor, a severability clause, alternative wording the circuit has already accepted. Abandonment is what 4.4 describes as the exception, the case where no redesign closes the gap between the regulation's purpose and its statutory hook, not the default response to a high score. Second, where abandonment is the honest answer, JIA has not created that choice; it has surfaced it earlier. A municipality regulating in a genuinely unauthorized way was always going to lose that authority eventually, either through JIA's diagnosis before adoption or through litigation after it. The chilling effect this objection worries about is real only if JIA's diagnosis is wrong more often than a court's eventual ruling would be — and nothing about running the analysis earlier makes the underlying legal question harder to answer correctly. What changes is only when the municipality learns the answer, not what the answer is.

4.5 Governance and Implementation Choices

Two institutional alternatives to running JIA in-house deserve mention, each rejected for a specific reason. Retaining outside counsel for every ordinance in a municipality's annual docket was considered and set aside on cost grounds, not quality grounds. The expert judgment Part V's case studies show is real and valuable, but sustaining it across a full docket is not a cost most municipalities, particularly the small ones Part VI scales for, can carry. It also says nothing about which ordinances in that docket warrant the expense in the first place. A mandatory state-level pre-screening body was considered and set aside for a different reason: it would recreate, at the state level, the mismatch Part III found in a *wetgevingstoets*-style test. The risk this Article addresses turns on each municipality's own circuit, its own enabling statute, and its own fiscal position, none of which a single state office can screen for uniformly. It would also sit uneasily beside the home-rule structures Part I discusses, requiring state review of decisions home rule assigns to the municipality in the first place.

A related governance question arises once a score exists: what happens when MLVI returns high risk and the council wants to adopt the ordinance anyway. JIA does not resolve that question by design, and should not. The Protocol produces a documented score and a set of redesign options; the decision to proceed against that score, accept a specific redesign, or decline to adopt at all remains the council's, exercised on the public record the score itself helps create. The framework therefore preserves political accountability while improving legal transparency: the lawyer structures the information, the council still makes the choice. A high MLVI score is a reason to slow down and document the choice being made. It is not, and was never meant to be, a veto.

One further question is worth addressing directly: why should this analysis fall to the municipality itself, rather than waiting for courts to develop more predictable doctrine, or for Congress to

legislate more clearly? Both would help, and neither is within a municipality's control. Courts resolve the cases in front of them; they do not audit a city's regulatory docket in advance. Congress amends statutes at its own pace, on its own priorities, not a given municipality's timeline. JIA is addressed to the actor who can actually act before adoption — which is, by process of elimination, the municipality itself.

V. CASE STUDIES AND COMPARATIVE INSTITUTIONAL LESSONS

5.0 Selection Logic

These two cases were not chosen because they are representative of post-*Loper Bright* litigation generally. They were chosen because, together, they expose two different failure modes the framework is designed to catch. *Sackett* shows an agency's interpretation of a jurisdictional term failing against a private landowner; *San Francisco* shows an agency's interpretation of a permit-conditioning power failing against a municipality. One concerns the scope of a definition; the other concerns the scope of a delegated power to impose conditions. Different forms of statutory ambiguity, different parties on the losing end of the agency's argument, the same underlying dynamic: an agency's reading of its own authority not surviving independent judicial construction. Both cases also arise under the Clean Water Act, and that is a genuine limit on what they can show, not an incidental detail to set aside. Neither case tests JIA against a constitutional claim, an equal-protection or procedural-due-process challenge, a dispute over contested facts rather than contested text, or a question of purely discretionary enforcement. Block 1's ambiguity scan is built for statutory interpretation specifically; nothing in this Article suggests it extends to those other categories, and a municipality facing one of them should look elsewhere for guidance. What these two cases can show, and are used here to show, is how the framework performs against the specific

problem it was built to diagnose: an agency's or a regulator's reading of ambiguous statutory text meeting independent judicial review for the first time.

5.0.1 Why Retrospective Case Analysis

Testing JIA against decided cases carries a known limitation worth naming rather than setting aside. Knowing that *Sackett* and *San Francisco* reached the Supreme Court, and knowing how each came out, makes it easier to say in retrospect that Block 1 would have flagged the relevant language than it would have been to say so before either case was famous enough to write an article about. A prospective test — scoring a representative sample of currently pending ordinances against MLVI and tracking which ones eventually draw litigation — would answer the reproducibility and predictive-value questions this Article cannot fully answer with two already-decided cases. That test is exactly what Part VI's empirical agenda proposes. Retrospective analysis here does something narrower and more modest: it shows that the specific textual and institutional features Block 1 and Block 2 are built to flag were genuinely present in both cases, on the record, before either was decided — not that the framework would have picked these two cases out of a large pool with no case-specific knowledge at all. The stronger claim waits on the empirical work; this Article makes only the narrower one.

5.1 Retrospective Diagnosis: *Sackett v. EPA* and the Cost of an Undiagnosed WOTUS Provision

Sackett predates *Loper Bright* by a year. No municipality in 2023 could have applied a framework like JIA in real time, since the Protocol didn't yet exist. What the case offers instead is a controlled test of the diagnosis: had a municipality carried WOTUS-dependent permitting language on its books that year, would Block 1 and Block 2 have flagged the instability before the Supreme Court did?

Block 1 would have caught it without much trouble. The term “waters of the United States” had been reversed and re-reversed by the agencies for over a decade. Its “significant nexus” gloss traced back to a plurality opinion in *Rapanos v. United States*, not a majority — borrowed, unsettled language of exactly the kind Block 1 exists to catch. Block 2 would have pushed the score well outside the low-risk tier on that basis alone, even without a full accounting of D2, D4, and D5. D1 scores high because the term already carried a fractured judicial history. D3 scores high because WOTUS jurisdiction had produced circuit splits and repeat Supreme Court intervention since 2006 — a pattern that predicts volatility on its own, with no need to guess how *Sackett* in particular would come out.

A well-run scan wouldn’t have foreseen the “continuous surface connection” test *Sackett* eventually adopted. No ex-ante analysis could have named that specific rule in advance. What it would have shown is simpler: the term’s instability was already visible, on the record, long before the Court acted on it.

The unanimity of *Sackett*’s judgment hides how unsettled the underlying term still is. All nine justices agreed the Sacketts’ land fell outside the Clean Water Act. Only five joined the reasoning that got there.⁷⁸ Justice Kavanaugh, joined by Sotomayor, Kagan, and Jackson, concurred in the judgment while rejecting the majority’s “continuous surface connection” test as narrower than the statute’s actual word — “adjacent” — warranted, pointing to forty-five years of consistent agency practice reading that term more broadly. Justice Kagan, in a separate concurrence joined by the same three justices, went further, accusing the majority of substituting its own environmental

⁷⁸ *Sackett v. EPA*, 598 U.S. 651, 662 (2023) (majority opinion, Alito, J.); *id.* (Kavanaugh, J., concurring in the judgment, joined by Sotomayor, Kagan, and Jackson, JJ.) (arguing “adjacent” wetlands should include those separated from covered waters by a berm, dune, or levee); *id.* (Kagan, J., concurring in the judgment, joined by Sotomayor and Jackson, JJ.); *Rapanos v. United States*, 547 U.S. 715 (2006) (plurality opinion); *see supra* Part I, note 6.

policy preferences for the text Congress wrote. Four of nine justices would have drawn the jurisdictional line to include wetlands separated from covered waters by a berm, dune, or levee — precisely the kind of feature a coastal municipality’s own ordinance is likely to reference. A municipality whose permitting language depends on where that boundary sits can’t treat *Sackett*’s unanimous judgment as settling the question Block 1 would flag. The judgment resolved the Sacketts’ dispute. The reasoning behind it left the boundary that will decide the next one just as contested as it was the day before the decision came down.

Block 4’s redesign follows from that visible instability, not from any forecast of the specific rule to come. A municipality could have built its jurisdictional trigger around independent state wetlands authority, running alongside the federal WOTUS determination rather than depending on it. That reduces dependence on whatever the Court eventually does with the federal term, and it required no prediction of *Sackett*’s holding — only a willingness to treat a term with *Rapanos*’s track record as a standing liability, not settled law.

This is not an abstract concern about future litigation. The instability the concurrences flagged has already produced exactly the kind of state-by-state divergence Part III’s comparative discussion would predict from a system with thirteen independent interpreters and no single final word. The Army Corps paused approving jurisdictional determinations in *Sackett*’s immediate aftermath, just as it had after *Rapanos* nearly two decades earlier — the same borrowed, contested language producing the same administrative freeze twice. States moved in opposite directions once the freeze lifted: Illinois moved to legislate new wetlands protections to fill the gap *Sackett* opened, while neighboring Indiana rolled its own protections back the same year, and Missouri, North Carolina, and Tennessee followed Indiana’s direction while Colorado, New Mexico, and Washington followed Illinois’s. EPA did not issue its own implementation guidance on the

“continuous surface connection” test until March 2025, nearly two years after the decision.⁷⁹ A municipality drafting a wetlands ordinance anywhere in that two-year gap had no agency guidance to rely on and a live, unresolved disagreement among the justices about where the ordinance’s own key boundary sits. That is not a hypothetical Block 1 would flag. It is the condition Block 1 is built to catch.

5.2 *City and County of San Francisco v. EPA: A Validation, Not an Application*

San Francisco didn’t run JIA. The Protocol didn’t exist when the city filed its first challenge to the Oceanside permit in 2019. What the case offers instead is a real municipality, litigating for six years without JIA, arriving at the same moves Block 2 and Block 4 formalize — and winning with them. That shouldn’t surprise anyone. JIA formalizes reasoning experienced municipal litigators already apply under the pressure of a live dispute. Its contribution is timing: making that reasoning available before adoption, not after six years of litigation force it into view.

The city’s argument rested on a Block 1-style distinction. Section 1311(b)(1)(A)–(B) authorizes effluent limitations — concrete restrictions on a discharge. The “end-result” provisions the EPA added in 2019 were something else: they made the city responsible for the receiving water’s overall quality without specifying what conduct would satisfy that standard. Isolating that distinction, and treating the EPA’s broader Section 1311(b)(1)(C) authority as a separate question from its effluent-limitation authority, is exactly the move Block 1 asks a drafter to make before adoption. The Court’s holding followed the same line: specifying the steps a permittee must take is the EPA’s job, not an open-ended standard the permittee has to satisfy on its own.

⁷⁹ *Sackett Fallout Leaves Wetlands’ Fate to States*, E&E News (July 10, 2023) (describing the Army Corps’s pause on jurisdictional determinations following *Sackett* and a comparable pause following *Rapanos*); *Wetlands Protections After Sackett Decision*, EHSLeaders (2024) (describing divergent state legislative responses in Illinois and Indiana); EPA & U.S. Army Corps of Eng’rs, *Revised Definition of “Waters of the United States”*; *Response to Sackett v. EPA* (Mar. 2025) (implementation guidance issued approximately twenty-one months after the decision).

Justice Barrett’s dissent, joined by three other justices, took the opposite reading of the same statutory hook, and the disagreement is instructive for what JIA is meant to catch. Barrett argued that Section 1311(b)(1)(C)’s broad grant of authority to impose “any more stringent limitation” reaches conditions framed as results, not just specified conduct — she compared an end-result permit condition to an airline limiting checked-bag weight without dictating what a passenger packs. Separately, she warned that eliminating end-result conditions would force EPA to draft individualized, prescriptive terms for every permit, likely delaying or denying permits the agency doesn’t have the staff to particularize in time.⁸⁰

That warning was not speculative even at the time she wrote it. Massachusetts is one of four states with no delegated authority to administer its own NPDES program, leaving EPA to issue permits there directly — and as of October 2024, months before the decision came down, EPA already reported that roughly half of Massachusetts’s NPDES permit renewals were backlogged more than six months. The additional burden *San Francisco* now places on the agency — rewriting thousands of existing permits nationwide that relied on the end-result conditions the Court just eliminated — lands on a system that was already straining to keep pace before the decision added to its workload. The exposure is not confined to San Francisco’s own permit, either. Major cities with combined sewer systems — New York, Washington, Boston, and Buffalo among them — hold NPDES permits built around the same end-result structure the Court just eliminated, and some state regulators are already exploring whether to reintroduce receiving-water limitations through state law, echoing the same state-by-state divergence *Sackett*’s aftermath produced above.

⁸⁰ *City and County of San Francisco v. EPA*, 604 U.S. 334 (2025); *id.* (Barrett, J., dissenting in part, joined by Sotomayor, Kagan, and Jackson, JJ.) (arguing that Section 1311(b)(1)(C)’s “any more stringent limitation” language authorizes end-result conditions, and warning that eliminating them will delay or prevent permit issuance); Olympia Bowker, *San Francisco v. EPA: SCOTUS Limits Types of Clean Water Act Permit Conditions*, Anderson & Kreiger (Apr. 7, 2025) (reporting that, as of October 2024, EPA had a backlog of 49.2% of NPDES permit renewals in Massachusetts, one of four states without delegated NPDES authority); *see supra* Part I, notes 7-10.

That warning matters for a JIA analysis in a way the majority's holding alone doesn't. It's a live account of the administrative-cost side of the very redesign Block 4 recommends. A safe harbor or narrower compliance path isn't free to draft or administer. And a 5–4 decision, with four justices on the losing side, is itself a data point Block 2's D3 should weigh — not a settled question resolved unanimously, but a contested reading that could have gone the other way on a different panel. This illustrates the persistence of interpretive disagreement even at the Court's own highest level. It isn't the same phenomenon as the circuit-to-circuit divergence Part II calls the judicial lottery, but it's a reminder that even the Supreme Court's own resolution of a statutory hook can rest on a bare majority.

The city's own estimate of what six years of litigation were worth was ten billion dollars in avoided capital expenditure — the city's own advocacy figure, not an independently audited number. That estimate comes close to what a Block 3 forecast would have produced the day the end-result provisions were issued in 2019, long before six years and a Supreme Court argument stood between the city and that number.

Read together, the two cases share a structure rather than a lesson about municipalities specifically. In both, an agency's interpretation of a statutory hook failed to survive independent judicial construction — once against a private landowner, once against a city. The value of JIA lies less in producing a different legal analysis than in moving that analysis from litigation to drafting. Neither case shows JIA at work. Both show what JIA is for: reaching that conclusion before the litigation reaches it, not after.

5.2.1 Lessons Across the Cases

Read side by side, *Sackett* and *San Francisco* support four narrower conclusions than the cases individually might suggest. First, borrowed or long-contested statutory language is a more reliable

warning sign than any prediction about how a particular court will eventually rule on it — both cases turned on terms with histories of dispute measured in decades, not on language anyone could have flagged only in hindsight. Second, an agency’s own administrative record helps only once ambiguity has already been identified; in neither case did record quality substitute for a defensible reading of the statute itself. Third, fiscal exposure and doctrinal exposure move separately — *Sackett* carried real doctrinal instability with modest fiscal consequence for the government losing the argument, while *San Francisco*’s doctrinal question carried billions in fiscal consequence for the city that would have lost it. A framework that only measured doctrinal risk, or only measured cost, would have missed half of what each case shows. Fourth, in both cases the losing argument was available to be diagnosed well before the courts ruled on it, using tools that existed independently of the eventual outcome. That, more than either case’s result, is what this Article asks a reader to take from them.

These four points are not four independent observations so much as one point seen from four angles, and the angles are not accidental. Each map onto one of JIA’s own four blocks. The first is Block 1’s ambiguity scan, visible in the text before any court rules on it. The second is Block 2’s administrative-record dimension, necessary but never sufficient by itself. The third is the separation this Article has insisted on between Block 2’s doctrinal risk and Block 3’s fiscal forecast — two measurements the same underlying dispute can score very differently on. The fourth is the premise the whole framework rests on: that the diagnosis was available on the record before either case reached the Supreme Court, not read into the record afterward with the benefit of knowing how it came out. Two decided cases are not a validation study, and this Article has not claimed otherwise. But a four-block architecture, built independently of either case, mapping this cleanly

onto the axes that actually mattered in both is a stronger form of support than either case supplies standing alone.

The two cases point to something more specific than ambiguity itself, taken together. The municipality's exposure in both traced not to unclear drafting on its own part, but to dependence on a term or a power whose meaning another actor — an agency, ultimately a court — controlled and could revise. Before litigation, an agency supplies that meaning; once a challenge is filed, a court does. Ambiguity is the symptom Block 1 screens for. Dependence on someone else's contested reading is the condition underneath it. This Article calls that condition interpretive dependence: a regulation's reliance, for its own validity, on a reading of borrowed statutory language that the municipality does not control and cannot fix in place. Interpretive dependence, not ambiguity as such, is what Block 1 is ultimately screening for, and what the remaining blocks are built to price and redesign around. None of this means JIA predicts how a given panel will vote, or determines a statute's meaning in advance. *Sackett*'s own 5–4 split on the reasoning shows why no instrument could: competent, well-informed justices looking at the identical text still reached different readings. What JIA offers is the diagnosis that a term is contested enough to produce that kind of split, not a forecast of which side wins it.

5.3 A Note on Institutional Capacity: Ukraine's Decentralization Reform

The first two case studies validate JIA doctrinally, testing the framework against real litigation outcomes. The third serves a different function. It asks whether the broader institutional premise underlying JIA — that capacity built before an authority is withdrawn outperforms capacity assembled after — holds outside the American administrative-law setting that produced *Sackett* and *San Francisco* in the first place. The comparison is not between legal systems. It is between institutional responses to uncertainty from different sources: in the United States, uncertainty

follows independent judicial reinterpretation; in Ukraine, it followed the withdrawal of centralized administrative coordination. The mechanisms differ entirely. The institutional challenge — a local government suddenly responsible for decisions a higher authority had previously stabilized — does not.

Ukraine's post-2014 decentralization reform moved substantial fiscal and administrative authority down to newly amalgamated local governments. Thousands of small, largely insolvent local councils consolidated into approximately 1,469 territorial communities, or *gromadas*, between 2015 and 2020, each gaining direct fiscal transfers and a rising share of retained personal income tax.⁸¹ An OECD assessment of the reform's aftermath documents persistent, marked discrepancies in administrative, human-resource, and fiscal capacity across these communities⁸² — a finding consistent with, though narrower than, the institutional-readiness point this Article draws from it. Russia's full-scale invasion in February 2022 turned that peacetime observation into a live stress test. Martial law has produced two distinct governance arrangements across Ukraine's *gromadas*: in some, military administrations have replaced local self-government entirely; in others, elected local government continues, coordinating with regional and district authorities under wartime constraints. Field research in the hardest-hit regions — Chernihiv, Kherson, Kharkiv, and Zaporizhzhia — finds real continuity in local governance and service delivery despite the disruption. But the continuity is uneven in a way the peacetime data already predicted: some *gromadas*, including at least one oblast capital, established no dedicated response for displaced residents at all, while comparably situated communities built functioning housing and support

⁸¹ Valentyna Romanova & Andreas Umland, *Ukraine's Decentralization Reforms Since 2014: Initial Achievements and Future Challenges*, Chatham House Res. Paper (2019) (describing the pre-reform fragmentation of local government); see also OECD, *supra* note 80 (confirming the count of 1,469 resulting amalgamated municipalities).

⁸² OECD, *Rebuilding Ukraine by Reinforcing Regional and Municipal Governance* (2022) (documenting persistent discrepancies in municipal administrative, human-resource, and fiscal capacity).

programs within months of the invasion.⁸³ The variable separating the two outcomes was not the scope of authority each gromada held, which martial law had already constrained for all of them alike. It was the administrative capacity each had built before the war reached it.

The point isn't that OECD, or the wartime research that followed it, isolates the timing of capacity-building as the single decisive variable. It's that capacity, however acquired and whenever tested, is what differentiates communities operating under devolved authority — the only lesson this Article takes from the comparison.

Ukraine's judicial system, its civil-law tradition, its different relationship between local and central authority — none of that translates to the American setting, and this Article doesn't pretend otherwise. What travels is narrower: institutional resilience follows analytical capacity rather than legal authority alone, and in institutional settings of this kind, capacity built in advance costs less, and works better, than capacity assembled after the fact — a lesson that held under the ordinary strain of administrative reform, and has now held under the far greater strain of war.

The common thread running through all three examples is not environmental law, municipal law, or decentralization specifically. It is interpretive dependence, and what happens when the source of the interpretation it rests on is withdrawn or overruled. None of the three shows JIA, or anything like it, preventing the underlying event: the Supreme Court still narrowed “adjacent” in *Sackett*, the 1972 text of the Clean Water Act still constrained the EPA in *San Francisco*, and Russia's invasion still happened despite a decade of Ukrainian institution-building. What separated better from worse outcomes in each case was not the event itself, but whether the dependence it exposed

⁸³ Andrii Darkovich & Maryna Rabinovych, *Ukrainian Decentralisation Under Martial Law: A Balancing Act*, ZOIS Spotlight (Sept. 17, 2025) (describing the two governance arrangements produced by martial law and reporting field-research findings of continuity in local governance in war-affected regions); Valentyna Romanova, *Does Decentralization Boost Ukrainian Resilience? The Role of Local Authorities in Supporting Internally Displaced Persons*, 41 Post-Soviet Affairs 459 (2025) (documenting that some gromadas, including at least one oblast capital, established no dedicated response for internally displaced persons while comparably situated communities did).

had already been named, priced, and prepared for, or was instead discovered for the first time once the event arrived. Part VI turns from that observation to the practical question it raises: how a municipality without San Francisco's resources, or Ukraine's decade of preparation, actually builds that readiness in advance.

VI. IMPLEMENTATION ROADMAP FOR U.S. LOCAL GOVERNMENTS

6.0 Why Not Rely on Expert Judgment Alone

There's an obvious objection here, and it deserves a direct answer. San Francisco's litigation team, discussed in Part V, reached JIA's conclusions without JIA. If experienced municipal counsel already reason this way under pressure, what does formalizing that reasoning actually add?

Four things, and none of them are things good judgment alone can supply. Most municipalities don't have a city attorney's office that has litigated a case like San Francisco's, and no municipality can assume one into existence the next time an ordinance comes up for a vote. Intuitive risk assessment, however sound, doesn't reproduce the way a scored instrument does. Two lawyers reasoning informally about the same ordinance can land in different places without either of them being wrong, which is a shaky foundation for a decision a city council has to stand behind. Intuition also leaves nothing on paper, and paper is exactly what *Skidmore* now asks a court to weigh⁸⁴ — a conclusion an attorney reached in her head persuades no one on appeal. And intuition doesn't scale. A legal department fielding two hundred ordinances a year can't triage them by feel in any way it could later document, defend, or repeat. JIA doesn't replace the judgment San Francisco's counsel exercised. It makes that judgment available to municipalities that don't have counsel like that on staff, in a form that survives being written down.

⁸⁴ See *supra* Part II, note 25 (*Skidmore v. Swift & Co.*, 323 U.S. 134 (1944)).

6.1 Scaling by Municipal Size and Capacity

JIA only matters if a municipality without a large legal department can actually run it. The Protocol scales its demands with population and staffing rather than assuming a capacity most municipalities don't have. The population cutoffs and time estimates below are illustrative benchmarks drawn from the author's own Quick-Scan Checklist and its accompanying scaling model⁸⁵ — not tested thresholds. That's the same discipline applied to MLVI's own risk tiers in Part IV, and for the same reason: where these lines actually belong is a question for later, not a claim made here.

Large municipalities, above roughly 100,000 residents, can support a standing three-person group — the municipal attorney as lead, a finance or budget analyst, and a planning or regulatory specialist — convened whenever an ordinance scores 40 or above on the MLVI. Medium municipalities, between roughly 50,000 and 100,000, typically run JIA through a single attorney or a retained outside consultant, using a standardized checklist and short-report template that takes eight to twelve hours per ordinance and needs no specialized administrative-law training. Small municipalities, below roughly 50,000, usually have no separate budget analyst or planning specialist to call on at all. A single attorney using the Quick-Scan instrument can get through a preliminary review in two to four hours, and municipalities with no dedicated legal staff can pool resources through a regional JIA platform built around an existing Council of Governments,

⁸⁵ These implementation figures are illustrative estimates, not independently verified data. They reflect the author's own Municipal Quick-Scan Checklist and its accompanying scaling model, developed as a companion operational instrument to this Article's four-block methodology; see Oksana Manko, *Judicial Impact Assessment (JIA) Protocol: Municipal Quick-Scan Checklist* (unpublished manuscript, on file with author). Because that instrument is not yet independently published, a reader cannot verify these specific figures directly; they are offered here as a starting benchmark for the scaling logic Part VI describes, not as a validated timing study, and are themselves a candidate for the empirical testing Part VI proposes.

sharing one or two attorneys with administrative-law expertise across several small municipalities instead of each one trying to build that capacity alone.

None of this is a smaller municipality getting a weaker version of JIA. It's a design choice already built into Block 3 back in Part IV: where a municipality has no separate finance office to run the cost forecast, the same attorney who scores the MLVI runs that forecast too, in the same sitting, rather than the Protocol assuming a division of labor that doesn't exist on the ground.

6.2 Institutionalizing JIA: Integration into Drafting Procedure

A JIA sitting in a binder no one opens before a vote does no good. Embedding it means building an MLVI checkpoint into the process that already routes an ordinance through legal review before it reaches a council agenda, so scoring becomes a step in adoption itself rather than something someone has to remember to request.

Putting the score in the public record raises a question worth answering rather than waving off: does disclosing a high D1 or D3 score just hand a future plaintiff evidence the municipality knew about the risk and adopted the ordinance anyway? The concern is real, but the facts driving those scores are already public regardless of whether MLVI exists. Circuit precedent, the statute's own text, the state of the administrative record — none of that is a secret JIA creates. It's the same material a plaintiff's attorney would find doing the same homework the municipality just did first. What MLVI adds isn't new exposure. It's the municipality's own documented response to exposure that was already there. A municipality that scores a provision high risk and redesigns it before adoption builds a materially different file than one that scores the same provision high risk and does nothing. The first is documented, thorough reasoning of exactly the sort *Skidmore* instructs a court to evaluate — not a guarantee of winning, but a record capable of earning the persuasive weight *Skidmore* asks a court to consider. Municipalities still uneasy about this can

treat the underlying MLVI work product as privileged legal advice and place only the resulting drafting decision, not the full sub-scores, in the public record — though that protection is not automatic and depends on the specific documentation qualifying as material prepared in anticipation of litigation, not merely as part of ordinary legislative drafting.⁸⁶

Beyond any single municipality, groups like the National League of Cities, and other organizations positioned to play a similar role,⁸⁷ are positioned to do for JIA what they already do for other model practices — circulate standardized checklists, model ordinance language, and MLVI benchmarks that let smaller municipalities adopt the Protocol without building it from scratch. A state or local government law section of a bar association could develop a short certification curriculum for municipal attorneys along the same lines, giving JIA an institutional home instead of leaving it dependent on whichever attorney happens to have read this Article.

6.3 When to Reapply

JIA isn't a gate an ordinance passes through once and never sees again. A regulation scored low risk at adoption can turn into a higher-risk one later, and four events should trigger rescoring rather than waiting for the next lawsuit to force the question. A new Supreme Court or circuit decision touching a dimension the ordinance already leans on — usually D1 or D3. A change in the state's enabling statute or preemption posture, which hits D4. A material shift in the municipality's own fiscal position, which hits D5. And, for any ordinance stuck in the medium- or high-risk tier for a long stretch without redesign, a periodic look on a fixed schedule rather than an open-ended hold. None of this calls for rerunning all four blocks from scratch — only rescoring whatever dimension

⁸⁶ *Hickman v. Taylor*, 329 U.S. 495 (1947); Fed. R. Civ. P. 26(b)(3) (protecting materials “prepared in anticipation of litigation or for trial”). Whether routine, pre-adoption MLVI scoring satisfies this standard, as opposed to scoring conducted after a specific challenge is anticipated, has not been tested and may vary by circuit; municipalities relying on this option should confirm the point with counsel before treating MLVI work product as protected.

⁸⁷ See *supra* Part I, note 13 (National League of Cities model-practice materials). ICMA and similar practitioner associations plausibly serve a comparable function, though no specific ICMA publication is cited here.

the triggering event actually touches. Figure 2 places this reapplication cycle next to the adoption process itself, closing the loop Figure 1 in Part IV left open.




Figure 2. Reapplication cycle: rescoring is triggered by a defined event, not run on a fixed calendar alone.

6.4 An Empirical Agenda

Every provisional piece of MLVI flagged earlier in this Article — the equal weighting across five dimensions, the specific point ranges marking low, medium, and high risk, the assumption that a second reviewer would land near the same score working from the same record — is provisional because none of it has been tested against outcomes yet, not because testing it can't be done.

The next stage of this research program is that test. Score a representative sample of U.S. municipal ordinances against MLVI, track which ones actually drew litigation and which didn't, and use that to ask whether unequal weights predict outcomes better than the equal weighting used here, and whether Block 2's tier boundaries hold up against real cases or need to move. That work will also have to settle something this Article leaves open: what counts as a "litigation outcome" in the first place — a complaint filed, a judgment, an appeal, Supreme Court review, a settlement extracted under threat of suit. Different definitions will likely point to different answers about which dimensions actually predict exposure. Testing reproducibility directly — putting the same ordinance in front of independent attorneys and seeing whether their scores land close together — would answer the question this Article has so far only argued for by design, not shown with data. None of that work happens here. It's what a first-generation instrument is supposed to invite.

VII. CONCLUSION

Loper Bright did more than change how federal courts read statutes. It changed *when* a local government has to start thinking about judicial risk. Under *Chevron*, that moment came, if it came at all, after a challenge was filed, when an agency's reading of an ambiguous term still had a real shot at surviving on deference alone. That cushion is gone. A municipality's regulatory judgment now stands or falls on the same independent judicial construction traced through *Skidmore*'s uneven revival and the circuit-level divergence this Article calls the judicial lottery. The moment that matters has moved from the courtroom to the drafting table. Most municipalities have no process built for a moment that arrives that early.

Part III showed why the tools already in hand don't fill that gap. The Dutch *wetgevingstoets* answers a real question, just not this one — it was built for a legal system with no judicial lottery of the kind this Article identifies. American fiscal forecasting prices pension liabilities and debt

service with real sophistication, and treats litigation risk as an undifferentiated line in a general reserve fund, because no one ever asked it to do more. Neither tool was built with independent judicial reinterpretation of a specific ordinance's specific text in view. Pointing either one at the problem doesn't change that.

JIA is built with that view in mind. Four decisions, taken in sequence, each doing work the one before it makes possible. An ambiguity scan that treats borrowed statutory language as a standing liability rather than settled law. A scored litigation-risk profile that turns circuit precedent and administrative-record quality into a comparable, if not identical, assessment across attorneys. A fiscal forecast that prices a specific provision's exposure instead of letting it disappear into a general cushion. A redesign discipline — severability, safe harbors, layered drafting — that gives a drafter somewhere to go once the first three steps have told her where the risk actually sits. Part V tested that sequence against cases decided before the sequence existed. *Sackett* shows what the diagnosis would have caught. *San Francisco* shows a municipality reaching the same conclusion the hard way — six years and ten billion dollars after the fact. Neither case proves JIA works. Both show what it's for.

None of this settles into a finished instrument. The five MLVI dimensions are a starting point, not a validated model. The weights are equal because no better basis for unequal weights exists yet — not because equal weighting has been proven right. The risk tiers are illustrative benchmarks, not statistical cutoffs. Part VI laid out the empirical work that would turn a first-generation instrument into a tested one: score real ordinances, track real outcomes, check whether independent attorneys scoring the same record actually land in the same place. That work is the next stage of the same research program, not a footnote to it. This Article's job was never to finish it. Only to give it something concrete to test.

The literature on municipal power, discussed in Part I, has mostly asked a horizontal question: how much authority a city retains against the state that charters it. This Article has asked a vertical one: how much of that authority, once granted, survives a court's own independent reading of the statute it rests on. Both questions concern the same underlying condition of subnational governance — that local authority is borrowed, layered, and exercised at the sufferance of institutions a municipality does not control. Part V called this condition interpretive dependence. *Loper Bright* did not create it. It removed one of the buffers that had made it easier, for a while, to live with.

What this Article claims is narrower, and probably more durable for being narrower. The discipline *Loper Bright* now asks of courts — read the text, trace its structure, take the record seriously — is a discipline municipalities can practice too, before a court ever has to. A method built on that premise doesn't need to be finished to be useful. It needs only to be better than deciding by feel, and honest about where it might still be wrong. If *Loper Bright* has changed anything about what counts as competent regulatory drafting, it is this: interpretive dependence is no longer a background risk to manage after the fact, but a variable a careful drafter now has to price before the ordinance is written at all.